



{ name: mongo, type: DB }



mongoDB

MongoDB

The Complete Developer's Guide

Introduction to MongoDB

What is MongoDB?

MongoDB is a database which is created by the company who is also called MongoDB. The name stems from the word "humongous". This database is built to store a lot of data but also being able to work with the huge data efficiently. Ultimately, this is a database solution.

There are many database solutions such as MySQL, PostgreSQL, TSQL etc.

MongoDB is most importantly a database server that allows us to run different databases on it for example a Shop database. Within the database we would have different collections such as a Users collection or a Orders collection. We can have multiple databases and multiple collections per database.

Inside of the collection we have something called documents. Documents look like JavaScript JSON objects. Inside of a collection the documents are schema-less and can contain different data. This is the flexibility that MongoDB provides us with whereas SQL based database are very strict about the data stored within the database tables. Therefore, the MongoDB database can grow with the application needs. MongoDB is a No-SQL database.

Typically we will need some kind of structure in a collection because applications typically requires some type of structure to work with the data.

Diagram 1.1:



JSON (BSON) Data Format:

```
{
  "name": "Alex",
  "age": 29,
  "address": {
    "city": "Munich"
  },
  "hobbies": [
    { "name": "Cooking" },
    { "name": "Football" }
  ]
}
```



The above is an example of the JSON data format. A single document is surrounded by curly brackets. The data is normally structured with a Keys. Keys consist of a Name of the Key and a Key value. The Name of the Key (*which will be referred to as Key from now on*) and the Key Value must be wrapped around quotation marks (unless if the data is a type of number).

There are different types of values we can store such as: string, number, booleans and arrays.

We can also nest documents within documents. This allows us to create complex relations between

data and store them within one document, which makes working with the data and fetching data more efficient because it is contained in one document in a logical way. SQL in contrast requires more complex method of fetching data which require joins to find data in table A and data in table B to retrieve the relevant data.

Behind the scenes on the server, MongoDB converts the JSON data to a binary version of the data which can be stored and queried more efficiently. We do not need to concern ourselves with BSON as we would tend to work with JSON data.

The whole theme of MongoDB is flexibility, optimisation and usability and it is what really sets MongoDB apart from other database solutions because it is so efficient from a performance perspective as we can query data in the format we need it instead of running complex restructuring on the server.

The Key MongoDB Characteristics.

MongoDB is a no SQL solution because it is following an opposite concept/philosophy to SQL based databases. Instead of normalising the data i.e. storing data distributed across multiple tables where every table has a clear schema and then using relations, MongoDB goes for storing data together in a document. It does not force a schema hence schema-less/No-SQL.

We can have multiple documents in a single collection and they can have different structures as we have seen in Diagram 1.1. This is important, it can lead to messy data but it still our responsibility as developers to work with clean data and to implement a solution that works. On the other hand this provides us with a lot of flexibility. We could use MongoDB for applications that might still evolve, where the exact data requirements are not set yet. MongoDB allows us to started and we could always add data with more information in the same collection at a later point in time.

We also work with less relations. There are some relations, but with these embedded (nested) documents, we have less collections (tables) which we connect but instead we store data together. This is where the efficiencies is derived from, since data is stored together and when we fetch data from our application it does not require to reach out to multiple tables and merge the data because all the data is already within the single collection. This is where the speed, performance and flexibility comes from and can be seen beneficial for when building applications. This is the main reason why No-SQL solutions are so popular for read and write heavy applications.

MongoDB Ecosystem

The below Diagram 1.2 is the current snapshot of the MongoDB companies ecosystem and product offerings. The focus of this guide is on the MongoDB database used locally on our machines and on the cloud using Atlas. We will also dive in Compass and the Stitch world of MongoDB.

Diagram 1.2:



Installing MongoDB

MongoDB runs on all Operating Systems (OS) which include Windows/Mac/Linux. To install MongoDB we can visit their webpage on:

<https://www.mongodb.com/>

Under products select MongoDB server and download the MongoDB Community Server for our OS platform of choice. Install the MongoDB Server by following the installation steps.

Important Note: On Windows when installing click on the Custom Setup Type. MongoDB will be installed as a service which will be slightly different to how MongoDB runs on Mac & Linux.

On Mac and Linux we simply have a extracted folder which contains files. We would copy all the contents within this folder and paste them into any place within our OS i.e wherever we would want to install MongoDB.

We would then want to create a folder called data and a sub-folder called db anywhere within our OS, preferably in the root of the OS.

On Windows open up the command prompt or on Mac/Linux open up the terminal. This is where

we are going to spend most of our time using special commands and queries. Run the following command:

```
$ mongo
```

This should return command not found.

To fix this problem on a Mac go to the user folder and find a file called `.bash_profile` file (*if this does not exist we could simply create it*). Edit the file using a text editor. Add the following line:

```
export PATH=/Users/Username/mongobd/bin:$PATH
```

The path should be wherever we placed the MongoDB binary zip files. We need to add `:$PATH` at the end on Mac/Linux to make sure all our other commands work on our OS. Save the file and close the file.

Important Note: if you run into a problem on not being able to edit the `.bash_profile` using text editor use the following command to edit it within the terminal:

```
$ sudo nano ~/.bash_profile
```

This will allow you to edit the file within the terminal and enter the mongo bin file path. Press CTRL + o to save and CTRL + x to exit the nano edit.

To fix this on a Windows OS, we need to create an environment variable. Press the windows key and type environment which should suggest the Edit Environment Variable option. Under the user variables edit Path to add the directory path to where we installed the MongoDB files:

C:\Program Files\MongoDB\Server\4.0\bin

Restart the terminal/command prompt and now run the command:

\$ mongo

This should now return a error of connect failed on Mac/Linux.

On Windows it will connect because MongoDB is running as a service and has already started as a background service because we would have checked this during the installation. If we open the command prompt as administrator and ran the command 'net stop MongoDB' this will stop the background service running automatically and we can manually start and stop the MongoDB service running on windows. DO NOT RUN THIS COMMAND ON MAC/LINUX.

The mongo command is the client which allows us to connect to the server and then run commands on the databases. To start the service on Mac/Linux we would use the following command:

\$ mongod

When running this command to start the server it may fail if we chose a different default /data/db folder. If we used a different folder and not within the root of our OS we would need to start the mongod command instance followed by the --dbpath flag and the place where the /data/db is located if not within the root directory.

```
$ sudo mongod --dbpath "/data/db"
```

On Mac we would need to run the mongod command every time we wish to run the mongoDB service whereas on Windows this will run automatically even after restarting the system.

Now that we have the mongod server running minimise the terminal on Mac/Linux and open up a new terminal. We cannot close the mongod server terminal because it is running the service and if closed the mongoDB server everything will stop working and we cannot continue to work with the database server. Pressing the CTRL + C keys within the terminal will quit the mongod service, but we would need to re-run the mongod command again should we wish to run the server again.

We are now in the mongo shell which is the environment where we can run commands against our database server. We can create new databases, collections and documents which we will now focus on in the following sections.

Time to get Started

Now that we have the mongod server running and we can now connect to it using the mongo shell we can now enter the following basic commands in the mongo terminal:

Command	Description
<code>\$ cls</code>	Clear the terminal.
<code>\$ show dbs</code>	Display existing databases (there are three default databases: admin, config and local which store meta data).
<code>\$ use databaseName</code>	Connect/Switch to a database. If the database does not exist it will implicitly create a new database using the databaseName. It will not create the database until a collection and document is added.
<code>\$ db.collectionName.insertOne({"name of key": "key value"})</code>	Create a new collection. The db relates to the current connected database. This will implicitly create a new collection if it does not exist. We must pass at least one new data in the collection using the .insert() command passing in a JSON object. This will return the object to confirm the data was inserted into the database.

Important Note: we can omit the quotes around the name of the key within the shell but we must contain the quotes for the key value unless the key value is type of number. This is a feature within the mongo shell which work behind the scenes. MongoDB will also generate a uniqueId for new documents inserted into the collection.

Command	Description
<code>\$ db.collectionName.find()</code>	Display the document within the database collection.
<code>\$ db.collectionName.find().pretty()</code>	Display the documents within the database collection but prettify the data in a more humanly readable format.

This is a very basic introductory look at the following shell commands we can run in the mongo terminal to create a new database, switch to a database, create a collection and documents and display all the documents within a database collection either in the standard or pretty format.

Tip: to run the mongod server on a different port to the default port 27017 by run the following command. Note you would need to specify the port when running the mongo shell command as well. You would use this in case the default port is being used by something else.

```
$ sudo mongo --port 27018
```

```
$ mongo --port 27018
```

Shell vs Drivers

The shell is a great neutral ground for working with MongoDB. Drivers are packages we install for different programming languages the application might be written in. There are a whole host of drivers for the various application server languages such as PHP, node, C#, python etc. Drivers are the bridges between the programming language and the MongoDB server.

As it turns out, in these drivers, we would use the same command as we use in the shell, they are just slightly adjusted to the syntax of the language we are working with.

The drivers can be found on the MongoDB website:

<https://docs.mongodb.com/ecosystem/drivers/>

Throughout this document we will continue to use the Shell commands as it is the neutral commands. We can take the knowledge of how to insert, configure inserts, query data, filter data, sort data and many more shell commands. These commands will continue to work when we use the drivers but we would need to make reference to the driver documentation to understand how to use the shell commands but using the programming language syntax to perform the commands using the drivers. This will make us more flexible with the language we use when building applications that uses MongoDB.

MongoDB & Clients: The Big Picture

Diagram 1.3:



MongoDB & Clients: The Big Picture

Diagram 1.4:



As we can see in Diagram 1.3 the application driver/shell communicates to the MongoDB server. The MongoDB server communicates with the storage engine. It is the Storage Engine which deals with the data passed along by the MongoDB Server, and as Diagram 1.4 depicts it will read/write to database and/or memory.

Understanding the Basics & CRUD Operations

Create, Read, Update & Delete (CRUD)

We could use MongoDB to create a variety of things such as an application, Analytics/BI Tools or data administration. In an application case, we may have an app where the user interacts with our code (the code can be written in any programming language) and the mongoDB driver will be included in the application. In the case of a Analytics/BI Tools we may use the BI Connector/Shell provided by mongoDB or another import mechanism provided by our BI tool. Finally, in the database administrator case we would interact with the mongoDB shell.

In all the above cases we would want to interact with the mongoDB server. In an application we would typically want to be able to create, read, update or delete elements e.g. a blog post app. With analytics, we would at least want to be able to read the data and as an admins we would probably want to do all the CRUD actions.

CRUD are the only actions we would want to perform with our data i.e. to create it, manage it or read it. We perform all these actions using the mongoDB server.

Diagram 1.5:

CREATE

```
insertOne(data, options)  
insertMany(data, options)
```

UPDATE

```
updateOne(filter, data, options)  
updateMany(filter, data, options)  
replaceOne(filter, data, options)
```

READ

```
findOne(filter, options)  
find(filter, options)
```

DELETE

```
deleteOne(filter, options)  
deleteMany(filter, options)
```

The above are the four CRUD operations and the commands we can run for each action. In later sections we will focus on each CRUD action individually to understand in-depth each of the actions and syntax/command we can use when performing CRUD operation with our mongoDB data collection and documents.

Understanding the Basics & CRUD Operations

Finding, Inserting, Updating & Deleting Elements

To show all the existing databases within the MongoDB server we use the command "show dbs" while we use the use followed by the database name to switch to a database. The db will then relate to the switched database.

To perform any CRUD operations, these commands must always be performed/executed on a collection where you want to create/update/delete documents. Below are example snippets of CRUD commands on a fictitious flights database (where the collection is called flightData).

```
$ db.flighData.insertOne( {distance: 1200} )
```

This will add a single document to the collection as we have seen previously.

```
$ db.flightData.deleteOne( {departureAirport: "LHR"} )
```

```
$ db.flightData.deleteMany( {departureAirport: "LHR"} )
```

The delete command takes in a filter. To add a filter we would use the curly brackets passing in

which name of key and key value of the data we wish to filter and delete. In the above example we used the `departureAirport` key and the value of `TXL`. The `deleteOne` command will find the first document in our database collection that meets the criteria and deletes it. The command will return:

```
{ "acknowledged" : true, "deleteCount" : 1 }
```

If a document was deleted in the collection this will show the number of deleted documents (the `deleteOne` command will always return 1). If no documents matched the filter and none were deleted the returned `deleteCount` value will be 0.

The `deleteMany` in contrast will delete many documents at once where the documents matches the filter criteria specified.

Note: The easiest way to delete all data in a collection is to delete the collection itself.

```
$ db.flightData.updateOne( {distance: 1200}, { $set: {marker: "delete"}} )
```

```
$ db.flightData.updateMany( {distance: 1200}, { $set: {marker: "delete"}} )
```

The `update` command takes in 3 argument/parameters. The first is the filter which is similar to the `delete` command. The second is how we would want to update/change the data. We must use the

`{ $set: { } }` keyword (anything with a \$ dollar sign in front of the keyword is a reserved word in *mongoDB*) which lets *mongoDB* know how we are describing the changes we want to make to a document. If the update key:value does not exist, this will create a new key:value property within the document else it will update the existing key:value with the new value passed in. The third parameter is options which we will analyse in great detail in the latter sections.

Important Note: when passing in a filter we can also pass in empty curly brackets `{ }` which will select all documents within the collection.

If successful with updating many this will return within the terminal an acknowledgement as seen below, where the number of matched the filter criteria and the number of data modified:

```
{ "acknowledged" : true, "matchedCount" : 2, "modified" : 2 }
```

If we were to delete all the documents within a collection and use the command to find data in that collection i.e using the `db.flightData.find().pretty()` command, the terminal will return empty/nothing as there are no existing documents to read/display.

The above demonstrates how we can find, insert, update and delete elements using the update and delete command.

Now we have seen how we can use `insertOne()` to add a single document into our collection. However, what if we want to add more than one document? We would use the `insertMany()` command instead.

```
db.flightData.insertMany( [  
  {  
    "departureAirport": "LHT",  
    "arrivalAirport": "TXL"  
  },  
  {  
    "departureAirport": "MUC",  
    "arrivalAirport": "SFO"  
  }  
])
```

We pass in an array of objects in order to add multiple documents into our database collection. The square brackets is used to declare an array. The curly brackets declare a object and we must use comma's to separate each object. If successful, this will return acknowledged of true and the `insertIds` of each object/document added into the collection.

Important Note: MongoDB by default will create a unique id for each new document which is assigned to a name of key called “_id” followed by a random generated key. When inserting a object we could assign our own unique id using the _id key followed by a unique value. If we insert a object and pass in our own _id key value and the value is not unique this will return a duplicate key error collection in the terminal. We must always use a unique id for our documents and if we do not specify a value for _id then MongoDB will generate one for us automatically.

Understanding the Basics & CRUD Operations

Diving Deeper Into Finding Data

Currently we have seen the .find() function used without passing any arguments for finding data within a collection. This will retrieve all the data within the collection. Just as we would use filter to specify a particular records or documents when deleting or updating a collection, we can also filter when finding data.

We can pass a document into the find function which will be treated as a filter as seen in the example below. This allows us to retrieve a subset of the data rather than the whole data within an application.

```
db.flightData.find( {intercontinental : true } ).pretty()
```

We can also use logical queries to retrieve more than one document within a collection that matches the criteria as demonstrated in the below example. We query using another object and then one of the special operators in MongoDB.

```
db.flightData.find( {distance: {$gt: 1000 } } ).pretty()
```

In the above we are using the \$gt: operator which is used for finding documents “greater than” the value specified. If we were to use the findOne() operator this will return the first record within the collection that matches the criteria.

Understanding the Basics & CRUD Operations

Update vs UpdateMany

Previously we have seen the updateOne() and updateMany() functions. However, we can also use another update function called update() as seen in the example below:

```
db.flightData.update( { _id: ObjectId("abc123") }, { $set: { delayed: true } } )
```

The update() function works exactly like the updateMany() function where all matching documents to the filter are updated. The difference between update() and updateMany() is that the \$set:

operator is not required for the update() function whereas this will cause an error for either the updateOne() and updateMany() functions. So we can write the above syntax like so and would not get an error:

```
db.flightData.update( { _id: ObjectId("abc123") }, { delayed: true } )
```

The second and main difference is that the update function takes the new update object and replaces the existing object (*this does not affect the unique id*) updating the document. It will only patch the update object instead of replacing the whole existing object (just like the updateOne() and updateMany() functions), if we were to use the \$set: operator, otherwise it would override the existing document.

This is something to be aware of when using the update() function. If we intend to replace the whole existing document with a new object then we can omit the \$set: operator. In general it is recommended to use updateOne() and updateMany() to avoid this issue.

If, however, we want to replace a document we should use the replaceOne() function. Again, we would place our filter and the object we want to replace with. This is a more explicit and more safer way of replacing the data in a collection.

```
db.flightData.replaceOne( { id: ObjectId("abc123") }, { departureAirport: "LHT", distance: 950 } )
```

Understanding the Basics & CRUD Operations

Understanding Find() & The Cursor Object

If we have a passengers collection which stores the name and age of passengers and we want to retrieve all the documents within the passenger collection we can use the find() function as we have seen previously.

```
db.passengers.find().pretty()
```

Useful Tip: when writing commands in the shell we can use the tab key to autocomplete for example if we wrote db.passe and tab on our keyboard, this should auto-complete db.passengers.

We will notice where a collection has many data, the find() function will return all the data but display all the data with the shell. If we scroll down to the last record we should see Type "it" for more within the shell. If we type the command it and press enter, this will display more data from the returned find() function. The find() command in general returns back what is called a Cursor Object and not all of the data.

The find() does not give an array of all the documents within a collection. This makes sense as the

collection could be really large and if the find() was to return the whole array, imagine if a collection had 2million documents – this could take a really long time but also send a lot of data over the connection.

The Cursor Object is an object that has many meta data behind it that allows us to cycle through the results, which is what the “it” command did. It used the Cursor Object to fetch the next group (cycle) of data from the collection.

We can use other methods on the find() function such as toArray() which will exhaust the cursor i.e. go through all of the cursors and fetch back all the documents within the array (i.e. not stopping after the first 20 documents – a feature within the mongoDB shell).

```
db.passengers.find().toArray()
```

There is a forEach method that can also be used on the find() function. The forEach allows us to write some code to do something on every element that is in the database. The syntax can be found within the driver documents for whichever language we are using for our application e.g. PHP or JavaScript etc. Below is a JavaScript function which the shell can also use:

```
db.passengers.find().forEach( (document) => { printjson(document) } )
```

The `forEach` function in JavaScript gets the document object passed in automatically into the arrow function and we can call this whatever we want i.e. `passengersData`, `data`, `x`, etc. In the above we called this `document`. We can then use this object and do whatever we want i.e. we used the `printjson()` command to print/output the document data as JSON. The above will also return all the documents within the collection because the `forEach` loops on every Cursor Object.

To conclude, the `find()` function does not provide us with all the documents in a collection even though it may look like it in some circumstances where there are very little data within a collection. Instead it returns a Cursor Object which we can cycle through the return more documents from the collection. It is unto us as the developer to use the cursor to either force it to get all the documents from a collection and place it in an array or better using the `forEach` or other methods to retrieve more than 20 documents (*the default number of items returned in the shell*) from the collection. Note the `forEach` is more efficient because it fetches/returns objects on demand through each iteration rather than fetching all the data in advance and loaded into memory which saves both on bandwidth and memory.

The Cursor Object is also the reason why we cannot use the `.pretty()` command on the `findOne()` function because the `findOne` returns one document and not a Cursor Object. For Insert, Update and Delete commands the Cursor Object does not exist because these methods do not fetch data, they simply manipulate the data instead.

Understanding the Basics & CRUD Operations

Understanding Projections



Imagine in our database we have the data for a person record and in within our application we do not need all the data from the document but only the name and age to display on our web application. We could fetch all the data and filter/manipulate the data within our application in any programming language. However, this approach will still have an impact on the bandwidth by fetching unnecessary data – something we want to prevent. It is better to filter the data out from the MongoDB server and this is exactly what projection allows us to do.

Below are examples of using projections to filter the necessary data to retrieve from our find query.

```
db.passengers.find( {}, {name: 1} ).pretty()
```

We need to pass in a first argument to filter the find search (note: a empty object will retrieve all documents). The second argument allows us to project. A projection is setup by passing another document but specifying which key:value pairs we want to retrieve back. The one means to include it in the data returned to us.

The above will return all the passengers document but only the name and id, omitting the age from the returned search results. The id is a special field in our data and by default it is always included. To exclude the id from the returned results, we must explicitly exclude it. To exclude something explicitly we would specify the name of key and set the value to zero as seen below:

```
db.passengers.find( {}, {name: 1, _id:0} ).pretty()
```

Note: we could do the same for age (e.g. age: 0), however, this is not required because the default is everything but the _id is not included in the projection unless explicitly specified using the one.

The data transformation/filtering is occurring on the mongoDB server before the data is shipped to us and is something that we would want because we do not want to retrieve unnecessary data which will impact on the bandwidth.

Understanding the Basics & CRUD Operations

Embedded Documents & Arrays

Embedded documents is a core feature of MongoDB. Embedded documents allows us to nest other documents within each other and having one overarching document in the collection.

There are two hard limits to nesting/embedded documents:

1. We can have up to 100 level of nesting (a hard limit) in MongoDB.
2. The overall document size has to be below 16mb

The size limit for documents may seem small but since we are only storing text and not files (we would use file storage for files), 16mb is more than enough.

Along with embedded documents, another documents we can store are arrays and this is not strictly linked to embedded documents, we can have arrays of embedded documents, but arrays can hold any data. This means we have list of data in a document.

Below are examples of embedded documents and arrays.

```
db.flightData.updateMany( {}, { $set: { status: { description: "on-time", lastUpdated: "1 hour ago" } } } )
```

In the above example we have added a new document property called status which has a embedded/nested document of description and lastUpdated. If we output the document using .find() function, the below document would now look something like the below:

```
{
  "_id": ...,
  "departure": "LHT",
  "arrivalAirport": "TXL",
  "status": {
    "description": "on-time",
    "lastUpdated": "1 hour ago",
  }
}
```

Note: we could add more nested child documents i.e. description could have a child nested document called details and that child could have further nested child documents and so on.


```
db.passengers.updateOne( {name: "Albert Twostone"}, {$set: {hobbies: ["Cooking", "Reading"]} } )
```

Arrays are marked with square brackets. Inside the array we can have any data, this could be multiple documents (i.e. using the curly brackets {}), numbers, strings, booleans etc.

If we were to output the document using the .find() function, the document would look like something below:

```
{  
  "_id": ...,  
  "name": "Albert Twostone",  
  "age": 63,  
  "hobbies": [  
    "Cooking",  
    "Reading"  
  ]  
}
```

Albert Twostone will be the only person with hobbies and this will be a list of data. It is important to note that hobbies is not a nested/embedded document but simply a list of data.

Understanding the Basics & CRUD Operations

Accessing Structured Data

To access structured data within a document we could use the following syntax:

```
db.passengers.findOne( {name: "Albert Twostone"} ).hobbies
```

We can specify the name of a structured data within a document by using the find query and then using the name of key we wish to access from the document, in the above we wanted to access the hobbies data which will return the hobbies array as the output:

```
["Cooking", "Reading"]
```

We can also search for all documents that have hobbies of Cooking using the syntax below as we have seen previously. This will return the whole document entry where someone has Cooking as a hobby. MongoDB is clever enough to look in arrays to find documents that match the criteria.

```
db.passengers( {hobbies: "Cooking"} ).pretty()
```

Below is an example of searching for objects (this includes searching within nested documents):

```
db.flightData.find( {"status.description": "on-time"} ).pretty()
```

We use the dot notation to drill into our embedded documents to query our data. It is important that we wrap the dot notation in quotations (e.g. "status.description") otherwise the find() function would fail.

This would return all documents (the whole document) where the drilled criteria matches. This allows us to query by nested/embedded documents. We can drill as far as we need to using the dot notation as seen in the example below:

```
db.flightData.find( {"status.details.responsible": "John Doe"} ).pretty()
```

This dot notation is a very important syntax to understand as we would use this a lot to query our data within our mongoDB database.

Understanding the Basics & CRUD Operations

Conclusion

We have now covered all the basic and core features of mongoDB to understand how mongoDB works and how we can work with it i.e. store, update, delete and read data within the database as well as how we can structure our data.

Understanding the Basics & CRUD Operations

Resetting The Database

To purge all the data within our MongoDB database server we would use the following command:

```
use databaseName
```

```
db.dropDatabase()
```

We must first switch to the database using the use command followed by the database name. Once we have switched to the desired database we can reference the current database using db and then call on the dropDatabase() command which will purge the specified database and its data.

Similarly, we could get rid of a single collection in a database using the following command:

```
db.myCollection.drop()
```

The myCollection should relate to the collection name.

These commands will allow us to clean our database server by removing the database/collections that we do not want to keep on our MongoDB server.

Schemas & Relations: How to Structure Documents

Why Do We Use Schemas?

There is one important question to ask – wasn't MongoDB all about having **no** data Schemas i.e. Schema-less. To answer this question, MongoDB enforces no Schemas. Documents do not have to use the same schema inside of one collection. Our documents can look like whatever we want it to look like and we can have totally different documents in one and the same collection i.e. we can mix different schemas.

Schemas are the structure of one document i.e. how does it look like, which fields does it have and what types of value do these fields have. MongoDB does not enforce schemas; however, that does not mean that we cannot use some kind of schema and in reality we would indeed have some form of schema for our documents. It is in our interest if we were to build a backend database that we have some form of structure to the types of documents we are storing. This would make it easier for us to query our database and get the relevant data and then cycle through this data using a programming language to display the relevant data within our application.

We are most likely to have some form of schemas because we as developers would want it and our applications will need it. Whilst we are not forced to have a schema we would probably end up with some kind of schema structure and this is important to understand.

Schemas & Relations: How to Structure Documents

Structuring Documents



We can use any of the structured approach in the diagram above depending on how we require it in our applications. In reality we would tend to use the approach in the middle or on the right.

The middle approach used the best of both worlds where there are some structure to the data, however, it also has the flexibility advantage that MongoDB provides us so that we can store extra information.

Note: we can assign the null value to properties in order to have a structured approach although the data may not have any actual values associated with the property. A null value is considered a valid value and therefore we can use a SQL (structured) type approach with all our documents.

There is no single best practice with how to set the structure of our data within our documents and it is up to us as developers to use the best structure that works best for our applications or whichever is to our personal preference.

Schemas & Relations: How to Structure Documents

Data Types

Now that we understand that we are free to define our own schemas/structure for our documents, we are now going to analyse the different data types we can use in MongoDB.

Data Types are the types of data we can save in the fields within our documents. The below table break the different data types for us:

Type	Example Value
String	"John Doe"
Boolean	TRUE
NumberInt (int32)	55, 100, 145
NumberLong (int64)	100000000000
NumberDecimal	12.99
ObjectId	ObjectId("123abc")
ISODate	ISODate("2019-02-09")
Timestamp	Timestamp(11421532)
Embedded Documents	{"a": {...}}
Arrays	{"b": [...]}

Notice how the text type requires quotation marks (*single or double*) around the value. There are no limitation in the size of the text. The only limitation is the 16mb for the whole document. The larger the text the larger the data it takes.

Notice how numbers and booleans do not require a quotation marks around the value.

There are different types of numbers in mongoDB. Integer (int32) are 32bit long numbers and if we try to store a number longer than this they would overflow that range and we will end up with a

different number. For longer integer numbers we would use NumberLong (int64). The integer solution we decide to choose will dictate how much space will be allocated and eaten up by the data. Finally, we can also store NumberDecimal i.e. numbers with decimal values (a.k.a float in other programming languages).

The default within the shell is to store a int64 floating point value but we also have a special type of NumberDecimal provided by MongoDB to store high precision floating point values. Normal floating point values (a.k.a doubles) are rounded and are not precise after their decimal place. However, for many use cases the floating point (double) is enough precision required e.g. shop store. If we are performing scientific calculations or something that requires a high precision calculation, we are able to use the special type that offers this very high decimal place precision (i.e. 34 digits after the decimal place).

The ObjectId is a special value that is automatically generated by MongoDB to provide a unique id but it also provides some temporal component that allows for sorting built into the ObjectId, respecting a timestamp.

The above table provides all the data types within MongoDB that we can use to store data within our database server.

Schemas & Relations: How to Structure Documents

Data Types & Limits

MongoDB has a couple of hard limits. The most important limitation: a single document in a collection (including all embedded documents it might have) must be less than or equal to 16mb. Additionally we may only have 100 levels of embedded documents.

We can read more on all the limitation (in great detail) on the below link:

<https://docs.mongodb.com/manual/reference/limits/>

For all the data types that mongoDB supports, we can find a detailed overview on the following link:

<https://docs.mongodb.com/manual/reference/bson-types/>

Important data type limits are:

- Normal Integers (int32) can hold a maximum value of +-2,147,483,674
- Long Integers (int64) can hold a maximum value of +-9,223,372,036,854,775,807
- Text can be as long as we want – the limit is the 16mb restriction for the overall document.

It's also important to understand the difference between int32 (NumberInt), int64 (NumberLong) and a normal number as you can enter it in the shell. The same goes for a normal double and NumberDecimal.

NumberInt creates a int32 value => NumberInt(55)

NumberLong creates a int64 value => NumberLong(7489729384792)

If we just use a number within the shell for example `insertOne({a: 1})`, this will get added as a normal double into the database. The reason for this is because the shell is based on JavaScript which only knows float/double values and does not differ between integers and floats.

NumberDecimal creates a high precision double value => NumberDecimal("12.99")

This can be helpful for cases where we need (many) exact decimal places for calculations.

When working with MongoDB drivers for our application's programming language (e.g. PHP, .NET, Node.js, Python, etc.), we can use the driver to create these specific numbers. We should always browse the API documents for the driver we are using within our applications to identify the methods for building int32, int64 etc.

Finally we can use the `db.stats()` command in the MongoDB shell to see stats of our database.

Schemas & Relations: How to Structure Documents

How to Derive Our Data Structure Requirements

Below are some guidelines to keep to mind when we think about how to structure our data:

- What data does our App need to generate? What is the business model?

User Information, Products Information, Orders etc. This will help define the fields we would need (and how they relate).

- Where do I need my data?

For example, if building a website do we need the data on the welcome page, products list page, orders page etc. This help define our required collections and field groupings.

- Which kind of data or information do we want to display?

For example the welcome page displays product names. This will help define which queries we need i.e. do we need a list of products or a single product.

These queries we plan to have also have an impact on our collections and document structure.

MongoDB embraces the idea of planning our data structure based on the way we retrieve the data so that we do not have to perform complex joins but we retrieve the data in the format or almost in

the format we need it in our application.

- How often do we fetch the data?

Do we fetch data on every page reload, every second or not that often? This will help define whether we should optimise for easy fetching of data.

- How often do we write or change the data?

Do we change or write data often or rarely will help define whether we should optimise for easy writing of data.

The above are things to keep in mind or to think about when structuring our data structures and schemas.

Schemas & Relations: How to Structure Documents

Understanding Relations

Typically we would have multiple collections for example a users collection, a product collection and a orders collections. If we have multiple collections that are relatable or where the documents in these relations are related, we obviously have to think about how do we store related data.

Do we use embedded documents because this is one way of reflecting a relation or alternatively, do we use references within our documents?

Nested/ Embedded Documents

Customers Collection

```
{
  "userName": "John",
  "age": 28,
  "address": {
    "street": "First Street",
    "City": "Chicago"
  }
}
```

References

Customers Collection:

```
{
  "userName": "Alan"
  "favBooks": ["id1", "id2"]
}
```

Books Collection:

```
{
  "_id": "id1",
  "name": "Lord of the Rings"
}
```

In the reference example above, we would have to run two queries to join the data from the different collections. However, if a book was to change, we would only update it in the books collection as the id would remain the same whereas in a embedded document relation we would have to update multiple customer records affected with the new change.

Schemas & Relations: How to Structure Documents

One to One Embedded Relation Example

Example:

One patient has one disease summary, a disease summary belongs to one patient.



Code snippet:

```
$ use hospital  
$ db.patients.insertOne( { name: "John Doe", age: "25", diseaseSummary: { diseases: ["cold",  
"sickness"] } } )
```

Where there is a strong one to one relation between two data, it is ideal to use a one to one embedded approach as demonstrated in the above example.

The advantage of the embedded nested approach is that within our application we only require a single find query to fetch the necessary data for the patient and disease data from our database collection.

Schemas & Relations: How to Structure Documents

One to One Reference Relation Example

Example:

One person has one car, a car belongs to one person.



Code snippet:

```
$ use.carData
```

```
$ db.persons.insertOne( { name: "John", age: 30, salary: 30000 } )
```

```
$ db.cars.insertOne( { model: "BMW", price: 25000, owner: ObjectId("5b98d4654d01c") } )
```

In most one to one relationships we would generally use the embedded document relations. However, we can opt to use a reference relation approach as we are not forced to use one approach.

For example, we have a more analytics use case rather than a web application and we have a use case where we are interested in analysing the person data and or analysing our car data but not so much in a relation. In this example we have a application driven reason for splitting the data.

Schemas & Relations: How to Structure Documents

One to Many Embedded Relation Example



Example:

One question thread has many answers, one answer belongs to one question thread.

Code snippet:

```
$ use support
```

```
$ db.questionThreads.insertOne( { creator: "John", question: "How old are you?", answers: [ { text: "I am 30." }, { text: "Same here." } ] } )
```

A scenario where we may use a embedded one to many relation would be post and comments. This is because you would often need to fetch the question along with the answers in an application perspective. Also usually there are not too many answers to worry about the 16mb document limit.

Schemas & Relations: How to Structure Documents

One to Many Reference Relation Example



Example:

One city has many citizens, one citizen belongs to one city.

Code snippet:

```
$ use cityData
```

```
$ db.cities.insertOne( { name: "New York City", coordinates: { lat: 2121, lng: 5233 } } )
```

```
$ db.citizens.insertMany( [ { name: "John Doe", cityId: ObjectId("5b98d6b44d") }, { name: "Bella Lorenz", cityId: ObjectId("5b98d6b44d") } ] )
```

In the above scenario we may have a database containing a collection of all major cities in the world and a list of every single person living within that city. It would seem to make sense to have a one to many embedded relationship, however, from an application prospective we may wish to only retrieve the city data only. Furthermore, a city like New York may have over 1 million people data and this would make fetching the data slow due to the volume of data passing through the wire. Furthermore, we may end up running into the document size limit of 16mb. In this type of scenario, it would make sense to split the data up and using the reference relation to link the data.

In the above we would only store the city metadata and will not store any citizen reference as this will also end up being a huge list of citizens unique id. Instead, we would create a citizens collection and within the citizens data we would make reference to the city reference. The reference can be anything but must be unique ie. we could use the ObjectId() or the city name etc.

This will ensure that we do not exceed the limitation of the 16mb per document as well as not retrieving unnecessary data if we are only interested in returning just the cities metadata from a collection.

Schemas & Relations: How to Structure Documents

Many to Many Embedded Relation Example



Example:

One customer has many products (via orders), a product belongs to many customers.

Code snippet:

```
$ use shop
```

```
$ db.products.insertOne( { title: "A Book", price: 12.99 } )
```

```
$ db.customers.insertOne( { name: "Cathy", age: 18, orders: [ { title: "A Book", price: 12.99, quantity: 2 } ] } )
```

We would normally model many to many relationships using references. However, it is possible to

use the embedded approach as seen above. We could store a collection for the products as meta data for an application to retrieve the data in order to help populate the embedded document of the customer collection using a programming language.

A disadvantage to the embedded approach is data duplication because we have the title and price of the product within the orders array as the customer can order the product multiple times as well as other customers which will cause a lot of duplication.

If we decide to change the data for the product, not only do we need to change it within the product collection but we also have to change it on all the orders affected by this change (or do we actually need to change old orders?). If we do not care about the product title changing and the price changing i.e. we have an application that takes a snapshot of the data, we may not worry too much about duplicating that data because we might not need to change it in all the places where we have the duplicated the data if the original data changes – this highly depends on the application we build. Therefore a embedded approach may work.

In other case scenarios where we absolutely need the latest data everywhere, a reference approach may be most appropriate in a many to many relationship. It is important to think about how we would fetch our data and how often do we want to change it and if we need to change it everywhere or are duplicate data fine before deciding which approach to adopt for many to many.

Schemas & Relations: How to Structure Documents

Many to Many Reference Relation Example



Example:

One book has many authors, an author belongs to many books.

Code snippet:

```
$ use bookRegistry
```

```
$ db.books.insert( { name: "favourite book", authors: [ objectId("5b98d9e4"), objectId("5b98d9a7") ] } )
```

```
$ db.authors.insertMany( [ { name: "Martin", age: 42 }, { name: "Robert", age: 56 } ] )
```

The above is an example of a many to many relation where a reference approach may be suitable for a scenario where the data that changes needs to be reflected everywhere else.

Schemas & Relations: How to Structure Documents

Summarising Relations

We have now explored the different relation options that are available to use. This should provide us enough knowledge to think about relations and when to use the most appropriate approach depending on:

- the application needs
- how often data changes
- if a snapshot data suffice
- how large is the data (how much data do we have).

Nested/Embedded Documents – group data together logically. This makes it easier when fetching the data. This approach is great for data that belong together and is not overlapping with other data. We should always avoid super-deep nesting (100+ levels) or extremely long arrays (16mb size limit per document).

References – split data across collections. This approach is great for related data but also shared data as well as for data which is used in relations and standalone. This allows us to overcome nesting and size limits (by creating new documents).

Schemas & Relations: How to Structure Documents

Using \$lookup for Merging Reference Relations

MongoDB has a useful operation called \$lookup that allows us to merge related documents that are split up using the reference approach.

The image on the right provides a scenario of a reference approach where the customer and books have been split into two



collections. The lookup operator is used as seen below. This uses the aggregate method which we have not currently learned.

```
$ customer.aggregate( [
  { $lookup: { from: "books", localField: "favBooks", foreignField: "_id", as: "favBookData" } }
])
```

The \$lookup operator allows us to fetch two related documents merged together in one document within one step (rather than having to perform two steps). This mitigates some of the disadvantages of splitting our documents across multiple collections because we can merge them in one go.

This uses the aggregate method framework (which we will dive into in later chapters) and within the aggregate we pass in an array because we can define multiple steps on how to aggregate the data. For now we are only interested in one step (a step is a document we pass into an array) where we pass the \$lookup step. The lookup passes in a document as a value, where we define 4 attributes:

- **from** – which other collection do we want to relate documents i.e. we would pass in the name of the collection where the other document lives that we wish to merge.
- **localField** – in the collection we are running the aggregate function on, where can the reference to the other (from) collection be found in i.e. the key that stores the reference.
- **foreignField** – which field are we relating to in our target collection (i.e. the from collection)
- **as** – provide an alias for the merged data. This will become the new key which the merged data will sit.

This is not an excuse to always using a reference relation approach because this costs more performance than having an embedded document.

If we have a references or want to use a references, we have the lookup step in the aggregate method that we can use to help get the data we need. This is a first look at aggregate and we will explore what else the aggregate can do for us in later chapters.

Schemas & Relations: How to Structure Documents

Understanding Schema Validation

MongoDB is very flexible i.e. we can have totally different schemas and documents in one and the same collection and that flexibility is a huge benefit. However, there are times where we would want to lock down this flexibility and require a strict schema.

Schema validation allows mongoDB to validate the incoming data based on the schema that we have defined and will either accept the incoming data for the write or update to the database or it will reject the incoming data and the database is not changed by the new data and the user gets an error.

validationLevel


Which document get validated?

Strict



All inserts & updates

Moderate



All inserts & updates
to correct documents

validationAction

What happens if validation fails?

Error



Throw error and deny
insert/update

Warn



Log warning but
proceed

Schemas & Relations: How to Structure Documents

Adding Collection Document Validation

To add schema validation in MongoDB and the easiest method is to add validation when we create a new collection for the very first time explicitly (not implicitly when we add a new data). We can use the `createCollection` to create and configure a new collection:

```
$ db.createCollection("posts", { validator: { $jsonSchema: { bsonType: "object", required: ["title",  
"text", "creator", "comments"], properties: { title: { bsonType: "string", description: "must be a string  
and is required." }, text: { bsonType: "string", description: "must be a string and is required" },  
creator: { bsonType: "objectId", description: must be an objectId and is required }, comments:  
{ bsonType: "array", description: "must be an array and is required", items: { bsonType: "object",  
required: ["text,"], properties: { text: { bsonType: "string", description: "must be a string and is  
required" }, author: { bsonType: "objectId", description: "must be an objectId and is required" }}}}} }  
}} )
```

The first argument to the `createCollection` method is the name of the collection i.e. we are defining the name of the collection. The second argument is a document where we would configure the new collection. The validator is an important piece of the configuration.

The validator key takes in another sub document where we can now define a schema against incoming data where inserts and updates has to validated. We do this by inserting a \$jsonSchema key with another nested sub document which will hold the schema.

We can add a son type with the value of object, so that everything that gets added to the collection should be a valid document or object. We can set a required key which has an array value. In this array we can define names of fields in the document which will be part of the collection that are absolutely required and if we try to add data that does not have these fields, we will get an error or warning depending on our settings.

We can add a properties key which is another nested document where we can define how for every property of every document that gets added to the collection will look like. In the example above we defined the title property, which is a required property, in more detail. We can set the bsonType which is the data type i.e. string, number, boolean, object, array etc. We can also set a description for the data property.

Because an array and has multiple items, we can add an items key and describe how the items should look like. We can nest this and this can have another nested required and properties keys for the items objects that exists within the array.

So the Keys to remember are:

The bsonType key is the data type.

The required key is an array of required properties that must be within an insert/update document.

The properties key defines the properties. This has sub key:value of of bsonType and description.

The Item key defines the array items. This can have sub key:value of all the above.

Important Note: it may be difficult to read in the terminal and may be easier to write in a text editor first and then paste into the terminal to execute the command. We can call the file validation.js to save the collection validation configuration. Visual Studio/Atom/Sublime or any other text editor/IDE will help with auto-formatting. Visual Studio has a option under code > Preference > Keyboard Shortcuts and then you can search for a shortcut command such as format document (shortcut is Shift + Option + F on a Mac).

We can now validate the incoming data when we explicitly create the new collection. We can copy the command from the text editor and paste it back into the shell and run the command to create the new collection with all our validation setup. This will return `{ "OK" : 1 }` in the shell if the new collection is successfully created.

If a new insert/update document fails the validation rules, the new document will not be added to the collection.

Schemas & Relations: How to Structure Documents

Changing the Validation Action

As a database administrator we can run the following command:

```
$ db.runCommand( { collMod: "post", validator: {...}, validationAction: "warn" } )
```

This allows us to run administrative commands in the shell. We pass a document with information about the command we wish to run. For example, in the above we run a command called collMod which stands for collection modifier, whereby we pass in the collection name and then we can pass in the validator along with the whole schema.

We can amend the validator as we like i.e. add or remove validations. In the above we added another administrative command after the validator document as a sibling called validationAction. The validationLevel controls whether all inserts and updates are checked or only updates to elements which were valid before. The validationAction on the other hand will either throw an "error" and stop the insert/update action or "warn" of the error but allow the insert/update to occur. The warn would have written a warning into our log file and the log file is stored on our system. We can update the validation action later using the runCommand() method as seen above.

Schemas & Relations: How to Structure Documents

Conclusion

Things to consider when modelling and structuring our Data.

- **In which format will we fetch your data?**

How does the application or data scientists need the data? We want to store the data in a way that it is easy to fetch especially in a use case where we would fetch a lot.

- **How often will we fetch and change the data?**

Do we need to optimise for writes or reads? It is often for reads but it may be different depending on the scenario. If we write a lot then we want to avoid duplicates. If we read a lot then maybe some duplicates are OK, provided these duplicates do not change often.

- **How much data will we save (and how big is it)?**

If the data is huge, maybe embedding is not the best choice.

- **How is the data related (one to one, one to many, many to many)?**

- **Will duplicate data hurt us (=> many Updates)?**

Do we update our data a lot in which we have to update a lot of duplicates. Do we have snapshot data where we do not care about updates to the most recent data.

- **Will we hit the MongoDB data/storage limit (embed 100 level deep and 16mb per document)?**

Modelling Schemas

- Schemas should be modelled based on application needs.
- Important factors are: read and write frequencies, relations, amount (and size) of data.

Schema Validation

- We can define rules to validate inserts and updates before writing to the database.
- Choose the validation level and action based on the application requirements.

Modelling Relations

- Two options: embedded documents or references.
- Use embedded documents if we have one-to-one or one-to-many relationships and there are no app or data size reasons to split the data.
- Use reference if data amount/size or app needs require it or for many-to-many relations.
- Exceptions are always possible — keep the app requirements in mind!

Useful Articles & Documents:

<https://docs.mongodb.com/manual/reference/limits/>

<https://docs.mongodb.com/manual/reference/bson-types/>

<https://docs.mongodb.com/manual/core/schema-validation/>

Exploring The Shell & The Server

Setting dbpath & logpath

In the terminal we can run the following command to see all the available options for our MongoDB server:

```
$ mongo --help
```

This command will provide a list of all the available options we can use to setup/configure our MongoDB server. For example the `--quiet` option allows us to change the way things get logged or output by the server.

Note: use the official document on the MongoDB website for more detailed explanation of all the available options.

The `--dbpath` arg and `--logpath` arg allows us to configure where the data and log files gets stored to because MongoDB writes our data to real files on our system. The logs allows us to see for example warnings of json schema validation as we seen in the last section.

We can create folders such as db and logs (these can be named as anything we want) and have

these folders located anywhere we want for example we could create it within the MongoDB directory which contains the bin folder and other related files.

If we start using mongod instance without any additional settings, it will use the root folder that has a data/db folder to store all our database records as a default setting. However, we can use the settings above to tell mongod to use another folder directory to store our data, the same is true for our logs.

When we start the instance of our MongoDB server, we can run the following command and passing in the options to declare the path of the dbpath and logpath as seen below:

Mac/Linux:

```
$ sudo mongod --dbpath /Users/userName/mongoDB/db
```

Windows command:

```
$ mongod --dbpath \Users\userName\mongoDB\db
```

Enter our password and this should bring up our MongoDB server as we have seen previously. We should now see in the db folder, MongoDB has created a bunch of files as it is now saving the data

in the specified folder that we passed into our command. This is now using a totally different database storage for writing all our data which is detached from the previous database storage of the default database path. Running the following command will also work for our logs:

Mac/Linux:

```
$ sudo mongod --dbpath /Users/userName/mongoDB/db --logpath /Users/userName/mongoDB/logs/logs.log
```

Windows command:

```
$ mongod --dbpath /Users/userName/mongoDB/db --logpath  
\Users\userName\mongoDB\logs\logs.log
```

The logs folder path requires a log file which we would define with a .log extension. This will automatically create and add a logs.log file within the directory path if the file does not exist when we run the command. All the output in the terminal will now be logged in the logs.log file compared to previously where it was logged in the terminal shell. This file can be reviewed for persistent and auditing of our server and viewing any warnings/errors.

This is how we set custom paths for our database and log files.

Exploring The Shell & The Server

Exploring the MongoDB Options

If we explore the different options in MongoDB using the `mongod --help` command in the terminal, there are many setup options available to us.

The WiredTiger options is related to our storage engine and we could either use the default settings or change some configurations if we know what we are doing.

We have useful commands such as `--repair` which we could run if we have any issues connecting or any warnings or issues related to our database files being corrupted. We could use the command `--directoryperdb` which will store each database in its own separate directory folder.

We could change the storage engine using the `--storageEngine` arg command, which by default is set to WiredTiger. Theoretically, MongoDB supports a variety of storage engines but WiredTiger is the default high performance storage engine. Unless, we know what we are doing and have a strong reason to change the engine, we should stick to the default.

There are other settings in regards to security which we will touch in the latter chapters.

Exploring The Shell & The Server

MongoDB as a Background Service

In the mongoDB options, there is an option called `--fork` which can only run on Mac and Linux.

```
$ mongod --fork --logpath /Users/userName/mongoDB/logs/logs.log/
```

The above fork command will error if we do not pass in a logpath to the log file. This command will start the mongoDB server as a child process. This does not block the terminal and we can continue to type in other commands in the same terminal with the server running. The server is now running as a background process instead of a foreground process which usually blocks the terminal window. In other words the mongoDB server is now running as a service (a service in the background). Therefore, in the same terminal we could run the mongo command to connect to the background mongoDB server service. This is also the reason why we require to pass in a logpath because the service is running in the background and it cannot log error/warning messages in the terminal, instead it will use/write the warning and errors in the log file.

On Windows, the fork command is unavailable. However, on Windows we can still startup mongoDB server as a service if we checked this option at the installation process. If we right click on command prompt and run as administrator, we can run the following command:

```
$ net start MongoDB
```

This will start up the MongoDB server as a background service. The question then becomes, how do we stop such a service?

On Mac we can stop the service by connecting to the server with the mongo shell and then switching to the admin database and running the shutdown server command to shut down the server we are connected to. Example commands below:

```
$ use admin
```

```
$ db.shutdownServer()
```

The exact same approach as the above will work on Windows. On Windows we also have an alternative method by opening the command prompt as administrator and running the following command:

```
$ net stop MongoDB
```

This is how we can use MongoDB server as a background service (instead of a foreground service) on either Mac, Linux or Windows.

Exploring The Shell & The Server

Using a Config File

Now that we have seen the various options we can set and use to run our MongoDB server, it is also worth noting that we can save our settings in a configuration file.

<https://docs.mongodb.com/manual/reference/configuration-options/>

This file could be automatically created for us when we run our MongoDB server, else we could create the config file ourselves and save this anywhere we want. We could create the config file within the Users/username/MongoDB/bin folder using a text editor such as VS Code to add the configuration code:

storage:

dbPath: "/Users/username/mongoDB/db/"

systemLog:

destination: file

path: "/Users/username/mongoDB/logs/logs.log/"

We can look at the documents or google search for more comprehensive config file setup.

Once we have the config file setup, how do we use the config file when we run an instance of the MongoDB server? MongoDB does not automatically pickup this file when we start to run the MongoDB server, instead when starting MongoDB we can use the following command to specify the config file the server should use:

```
$ sudo mongod --config /Users/userName/mongoDB/bin/mongod.cfg
```

```
$ sudo mongod -f /Users/userName/mongoDB/bin/mongod.cfg
```

Either above command will prompt MongoDB to use the config file from the path specified. This will start the MongoDB server with the settings setup in the configuration file. This is a useful feature because it allows us to save a snapshot of our settings (reusable blueprint) in a separate file which we can always use when starting up our MongoDB server. This also saves us time on writing a very long command prompt with all our settings when starting up our MongoDB server each time.

Important Note: we could use either .cfg or .conf as the file extension name when creating the MongoDB configuration file.

Exploring The Shell & The Server

Shell Options & Help

In this section we will go over the various shell options available to for us to use. Similar to the MongoDB server, there is a help option for the MongoDB shell:

```
$ mongo --help
```

This will provide all the command options for the shell. This has less options compared to the server because the shell is just a connecting client at the end of the day and not a server. We can use the shell without connecting to a database (if we just want to run javascript code) using the `--nodb` command, or we could use the `--quiet` command to have less output information in the terminal, we can define the port and host for the server using the `--port` arg and `--host` arg commands (by default it uses `localhost:27017`) and many more other options.

We can also add Authentication Options informations which we will learn more in later chapters.

In the shell we also have another command we can run:

```
$ help
```

This command will output a shortlist of some important help information/commands we can execute in the shell. We can also dive deeper into the help by running the help command followed by the command we want further help on, for example:

```
$ help admin
```

This will show further useful commands that we can execute when using the admin command e.g. admin.hostname() or admin.pwd() etc.

We can also have help displayed for a given database or collection in a database. For example:

```
$ use test
```

```
$ db.help()
```

We would now see all the commands that we did not see before that we can use on the new “test” database. We can also get help on the collection level which will provide a list of all the commands we can execute at the collection level.

```
$ db.testCollection.help()
```

Useful Links:

<https://docs.mongodb.com/manual/reference/configuration-options/>

<https://docs.mongodb.com/manual/reference/program/mongo/>

<https://docs.mongodb.com/manual/reference/program/mongod/>

Using the MongoDB Compass to Explore Data Visually

Exploring MongoDB Compass

We can download MongoDB Compass from the below link:

<https://www.mongodb.com/products/compass>

This is a GUI tool to interact with our MongoDB database. Once downloaded and installed on our machines we are ready to use the GUI tool. It is important to have the mongod server running in the background when we open the MongoDB Compass to connect to the database. We would connect to a Host and this by default will have localhost and port 27017. We can click connect and this will connect the GUI tool to the mongod server. We should be able to see the 3 default databases of admin, config and local.

We can now use the GUI tool to create a new database and collection name. Once a database and collection has been created we can then insert documents to the collection. We can also query our database documents.

We can now start using a GUI tool to interact with our database, collections and documents.

Note: it is best practice to learn how to use the shell first before using GUI tools.

Diving Into Create Operation

Understanding insert() Methods

We already understand that there are two methods for inserting documents into MongoDB which are `insertOne()` and `insertMany()` as an alternative. The most important thing to note is that `insertOne()` takes in a single document and we can but do not need to specify an id because we will get one automatically. The `insertMany()` does the same but with an array (list) of documents.

There is also a third alternative method for inserting documents called `insert()` – below is an example:

```
$ db.collectionName.insert()
```

This command is more flexible because it takes both a single document or an array of documents. `Insert` was used in the past but `insertOne` and `insertMany` was introduced on purpose so that we are more clear about what we will be inserting. Previously, in application code it was difficult to tell with the `insert` command whether the application was inserting a single or multiple documents and therefore may have been error prone.

There is also an importing data command as seen below:

```
$ mongoimport -d cars -c carsList --drop --jsonArray
```

The insert method can still be used in mongoDB but it is not recommended. The insert() method works with both a single document and multiple documents as seen in the examples below:

```
$ db.persons.insert( { name: "Annie", age: 20 } )
```

```
$ db.persons.insert( [ { name: "Barbara", age: 45 }, { name: "Carl", age: 65 } ] )
```

The output message in the terminal is also slightly different i.e. we would receive a text of:

```
$ WriteResult( { "nInserted" : 1 } )
```

```
$ BulkWriteResult( { "writeErrors": [], "writeConcernErrors": [], "nInserted": 2, "nUpserted": 0, "nMatched": 0, "nModified": 0, "nRemoved": 0, "upserted": [] } )
```

The above does not mean that the inserted document did not get an autogenerated id. The insert method will automatically create an ObjectId but will not display the ObjectId unlike the insertOne and insertMany commands output messages which does display the ObjectId. We can see the advantages of insertOne and InsertMany as the output message is a little more meaningful/helpful as we can immediately work with the document using the ObjectId provided (i.e. we do not need to query the database to get the new document id).

Diving Into Create Operation

Working With Ordered Inserts

When inserting documents we can define or specify some additional information. Lets look at an example of a hobbies collection where we keep track of all the hobbies people could possibly have when we insert many hobbies. Each hobby is a document with the name of the hobby:

```
$ db.hobbies.insertMany( [ { _id: "sports", "name": "Sports" }, { _id: "cooking", "name": "Cooking" },  
{ _id: "cars", "name": "Cars" } ] )
```

The id's for these hobbies can be auto-generated. However, there may be times when we want to use our own id because the data may have been fetched from some other database where we already have an existing id associated or maybe we need a shorter id. We can use `_id` and assign a value for the id. In the above the hobby name could act as a good id because each hobby will be unique. We must use `_id` and not just `id` if we want to set our own id for our documents.

Furthermore, the id must be unique else this would not work. We will no longer see an `ObjectId()` for these documents as we have used the `_id` as the unique identifier for the documents inserted.

If we try to insert a document with the same id we would receive an error message in the terminal referencing the index number (mongoDB uses zero indexing) of the document that failed the insert operation along with a description of duplicate key error.

```
$ db.hobbies.insertMany( [ { _id: "yoga", "name": "Yoga" }, { _id: "cooking", "name": "Cooking" }, { _id: "hiking", "name": "Hiking" } ] )
```

The above would fail due to the duplicate key error of cooking which was inserted previously in the above command. However, we would notice on the first item in the insertMany array i.e. Yogo will be inserted into the hobbies collectio, but the cooking and hiking documents will not be inserted into the collection due to the error. This is the default behaviour of mongoDB and this is called an ordered insert.

An ordered insert simply means that every element we insert is processed standalone, but if one fails, it cancels the entire insert operation but does not rollback the elements it has already inserted. This is important to note because it cancels the operation and does not continue to the next document (element i.e. hoking) which we would have known that it would have succeeded insert.

Often we would want this default behaviour, however, sometimes we do not. In these cases, we could override the behaviour. We would pass in a second argument, separated by a comma, to the insertMany command which is a document. This is a document that configures the insertMany operation.

```
$ db.hobbies.insertMany( [ { _id: "yoga", "name": "Yoga" }, { _id: "cooking", "name": "Cooking" }, { _id: "hiking", "name": "Hiking" } ], { ordered: false } )
```

The ordered option allows us to specify whether MongoDB should perform an ordered insert which is the default (*we could set this ordered option to true which is redundant because this is the default option*) or we could set this option to false which will make the insert operation not an ordered insert i.e. an unordered insert.

If we hit enter, we would still get a list of all the error, however, it will continue to the next document to perform the insert operation and this would insert the document that does not have any issues of duplicate keys i.e. hiking will now be inserted into the hobbies collection (yoga and cooking would fail due to the duplicate key issue).

By setting the ordered to false, we have changed the default behaviour and it is up to us to decide what we require or want in our application. It is important to note that this will not rollback the entire insert operation if something failed. This is something we will cover in the Transactions chapter. We can control whether the operation continues with the other documents and tries to insert everything that is perfectly fine.

We may use an unordered insert where we do not have much control with what is inserted into the database but we do not care about any document that fail because they already exist in the database. We could add everything that is not in the database.

Diving Into Create Operation

Understanding the writeConcern

There is a second option we can specify on `insertOne` and `insertMany` which is the `writeConcern` option. We have a client (either the shell or the application using a MongoDB server) and we have our MongoDB server. If we wanted to insert one document in our MongoDB server, on the MongoDB server we have a so called storage engine which is responsible for really writing our data onto the disk and also for managing it in memory. So our write might first end up in memory and there it manages the data which it needs to access with high frequency because memory is faster than working with the disk. The write is also scheduled to then end up on the disk, so it will eventually store data on the disk. This is true for all write operations i.e. `insertMany` and `update`.

We can configure a so-called `writeConcern` on all the write operations with an additional argument, the `writeConcern` which is another document where we can set settings.



The `w:` (default) option tells the MongoDB server of how many instances we want the write to be acknowledged. The `j:` option stands for journal which is an additional file which the storage engine manages, which is like a To-Do file. The journal can be kept to then for example perform save operations that the storage engine needs to do but have not been completed yet.

The storage engine is aware of the write and that it needs to store the data on disk just by having the write being acknowledged and being in memory. The idea behind a journal file is to make the storage engine aware of this and if the server should go down for some reason or anything else should happen, the journal file is there. If we restart the server or if the server recovers, the server can look to this file and see what it needs to do. This is a nice backup because the memory might have been wiped by then. The journal acts as a backup to-do list for the storage engine.

Writing into the database files is more performance heavy whereas a journal is like a single line which describes the write operations. Writing into the database is of course a more complex task because we need to find the correct position to insert the data and if we have indexes we also need to update these as well and therefore takes longer to perform. Writing in a to-do type list is much quicker.

We can set the `j: true` as an option which will now report a success for a write operation when it has been acknowledged and has been saved to the journal. This will provide a greater security.

There is a third option to the writeConcern which is the wtimeout: option. This simply sets the timeframe that we give out server to report a success for the write before we cancel it. For example, if we have some issues with the server connection or anything of that nature, we may simply timeout.

If we set the timeout value to a very low number, we may get more fails even though there is no actual problem, just some small latency.

```
{ w:1, j: undefined }
```

```
{ w:1, j: true }
```

```
{ w:1, timeout: 200, j: true }
```

This is the writeConcern option we can add to our write operations and how we can control this using the above document settings. Enabling the journal would mean that our writes will take longer because we do not only tell the server about the write operation but we also need to wait for the server to store the write operation in the journal, however, we get higher security that the write also succeeded. These option will again depend on our application needs.

Diving Into Create Operation

The writeConcern in Practice

Below is an example of using the writeConcern:

```
$ db.persons.insertOne( { name: "Alan", age: 44 }, { writeConcern: { w: 1, j: true } } )
```

The w:1 (default) simply means to make sure the server acknowledged the write operation. Note we could set this value to 0 which will return {"acknowledged": false} in the terminal when we insert the document. This option sends the request and immediately returns without waiting for a response of the request from the server. The storage engine had no chance of storing it in memory and to generate an objectId, hence why we receive {"acknowledged": false} in the terminal. This makes the write super fast because we do not have to wait for any response but we do not know where the write succeeded or not.

The journal by default is set to undefined or false. We can set this option to j: true. The output in the terminal does not change. The write will be slightly slower (*note if playing around locally we would not notice any change in speed*) because the engine would add the write to the journal and we would have to wait for that to finish before the operation is completed. This will provide a higher security by ensuring the write appears in the to-do list of the storage engine which will eventually

lead to the write operation occurring on the database.

Finally, the `wtimeout:` option is used to set a timeout operation. This allows us to set a time frame for the write operation so in the case where within a certain period within the year we would have shaky connections, we would rather have the write operation fail and we recognise it in our client application (we would have access to the error) and therefore try again at a later time without having to wait unnecessarily for the write operation where we have shaky connections.

Diving Into Create Operation

What is Atomicity?

Atomicity is a very important concept to any write operation. Most of the time the write operation e.g. `InsertOne()` would succeed, however, it can fail (there can be an error). These are errors that occur whilst the document is being inserted/written to memory (i.e. whilst being handled by the storage engine). For example, we were writing a document for a person including name, age and an array of hobbies, the name and age were written but then the server had issues and was not able to write the hobbies to memory. MongoDB protects us against this as it guarantees us an atomic transaction. This means the transaction either succeeds as a whole or it fails as a whole. If it fails during the write, everything is rolled back for the document we inserted.

This is important as it is on a per document level. The document means the top level document, so it includes all embedded documents and all arrays.

If we insertMany() where there are multiple documents being inserted into the database and the server fails during a write, we do not get atomicity because it only works at a document level. If we have multiple documents in one operation like the insertMany() operation, then only each document on its own is guaranteed to either fail or succeed but not on insertMany. Therefore, if we have issues during the insertMany operation, only the documents that failed are not inserted and then the exact behaviour will depend on whether we used ordered or unordered inserts but the document already inserted will not be rolled back.

We are able to control this on a bulk insert or bulk update level using a concept called transactions which we will look at in a later section as it requires some additional knowledge about MongoDB and how the service works.



MongoDB CRUD Operations are Atomic on the document level (including Embedded documents)

Diving Into Create Operation

Importing Data

To import data into our database, we must first exit the shell by pressing the control + c keys on our keyboard.

In the normal terminal, we need to navigate to a folder that contains the JSON file that we would want to import (JSON files can be imported) using the cd command. We can use the ls command to view the list of items within the directory we are currently in.

Once navigated to the folder containing the import file, we can run the following command:

```
$ mongoimport tv-shows.json -d moviesData -c movies --jsonArray --drop
```

The mongoimport command should be globally available since we added the path to our mongo binary to our path variables on our operating systems. If we did not do this, we need to navigate into the folder where our MongoDB binaries are in order to execute the mongoimport command above.

The first argument we pass is the name of the file we want to import (if we are not in the path of the located file we would have to specify the full folder path along with the file name). We then specify the database we want to import the data into using the -d flag. We can also specify the collection by using the -c flag.

If the JSON file holds array of documents we must also specify the --jsonArray flag to make the mongoimport command aware of this fact about the import data.

The last argument option we can add to the import command is the --drop which will tell the mongoimport that should this collection should already exist, it will be dropped and then re-added, otherwise it will append the data to the existing collection.

Important Note: the mongod server should be running in the background when we use the import command. When we press enter to execute the command, this will return in the terminal the connected to: localhost, dropping: moviesData.movies and imported: # documents in the terminal as a response to inform us which MongoDB server it is connected to, whether a collection was dropped/deleted from the database collection and the total number of data imported into the database collection.

Methods, Filters & Operators

In the shell, we access the database with the `db` command (this will differ slightly in a MongoDB drive). We would get access to a database and then to a collection in the database. Now we can execute a method like `find`, `insert`, `update` or `delete` on the collection. We would pass some data into the method as parameters/arguments for the method. These are usually a key:value pair where one is the field and the other is the value for that field name (documents are all about field and values or key and values).



The argument in the above example happens to also be a filter because the find method accepts a filter. It can use a filter to narrow down the set of documents it returns to us. In the above we have a equality or single value filter where the data is exactly the criteria i.e. equality.

We can also use more complex filters as seen in the below example. We have a document which has a field and its value is another document which has an operator as a field followed by a value.



We can recognise operators by the dollar sign \$ at the beginning of the operator. These are all reserved fields which are understood by mongodb. The operator in the above example is called a range filter because it does not just filter for equality, instead this will look for all documents that have an age that is greater than (\$gt) the value i.e. 30.

this is how the Read operator works and we will look at various different operators and the different ways of using them and the different ways of filtering the data that is returned to us. This is the structure we should familiarise ourselves with for all of our Read operations.

Diving Into Read Operation

Operators and Overview

There are different operators that we can differentiate into two groups:

- Query Selectors
- Projection Operators

Query selectors such as `$gt` allows us to narrow down the set of documents we retrieve while projection operators allows us to transform/change the data we get back to some extent. Both the Query and Projection operators are Read related operators.

Aggregation allows us to read from a database but also perform more complex transforms. This concept allows us to setup pipeline of stages to funnel our data through and we have a few operators that allow us to shape the data we get back to the form we need in our application.

- Pipeline Stages
- Pipeline Operators

The Update has operators for the fields and arrays. Inserts have no operators and deletes uses the same operators as the Read operators.

How do operators impact our data?

Type	Purpose	Change Data?	Example
Query Operator	Locate Data	No	\$eq \$gt
Project Operator	Modify data Presentation	No	\$
Update Operator	Modify & Add additional data	Yes	\$inc

Diving Into Read Operation

Query Selectors & Projection Operators

There are a couple of categories for Query Selectors:

Comparison, Logical, Element, Evaluation, Array, Comments & Geospatial.

Projections Operators we have:

\$, \$elemMatch, \$meta & \$slice

Diving Into Read Operation

Understanding findOne() and find()

The findOne() method finds exactly one document. We are able to pass in a filter into the method to define which one document to return back. This will find the first matching document.

```
$ db.movies.findOne( )
```

```
$ db.movies.findOne( { } )
```

Both of the above syntax will return the first document within the database collection. Note, this does not return a cursor as it only returns one document.

The alternative to findOne() is the find() method. The find() method will return back a cursor. This method theoretically returns one document, but since it provides us a cursor, it does not give us all the document but the first 20 documents within the shell.

```
$ db.movies.find( )
```

To narrow the find search we would need to provide a filter. To provide a filter we would pass in a document as the first argument (this is true for both find and findOne methods). The difference would be that findOne will return the first document that meets the criteria while the find method

will return all documents that meet the criteria.

```
$ db.movies.findOne( { name: "The Last Ship" } )
```

```
$ db.movies.findOne( { runtime: 60 } )
```

```
$ db.movies.find( { name: "The Last Ship" } )
```

```
$ db.movies.find( { runtime: 60 } )
```

To filter the data, we would specify the name of the field/key followed by the value we are expecting to filter the field by. In the above example we are filtering the name of the movie to be "The Last Ship". By default MongoDB will try to find the filter by equality.

This is the difference between find and findOne and how we would pass in a filter to narrow down the return read results. It is important to note that there are way more operators and ways to filtering our queries to narrow down our Read results when using either of the find commands.

Diving Into Read Operation

Working with Comparison Operators

In the official documentation we can view all the various operations available to us:

<https://docs.mongodb.com/manual/reference/operator/query/>

We will explore some of the comparative operators in the below examples:

```
$ db.movies.find( { runtime: 60 } ).pretty( )
```

```
$ db.movies.find( { runtime: { $eq: 60 } } ).pretty( )
```

The \$eq operator is the exact same as the default equality query which will find the document that matches equally to the query value which in the above case is runtime = 60.

```
$ db.movies.find( { runtime: { $ne: 60 } } ).pretty( )
```

This will return all documents that have a runtime not equal to 60.

```
$ db.movies.find( { runtime: { $gt: 40 } } ).pretty( )
```

```
$ db.movies.find( { runtime: { $gte: 40 } } ).pretty( )
```

The \$gt return all documents that have a runtime greater than to 40 while the \$gte returns greater than or equal to 40.

```
$ db.movies.find( { runtime: { $lt: 40 } } ).pretty( )
```

```
$ db.movies.find( { runtime: { $lte: 40 } } ).pretty( )
```

The \$lt return all documents that have a runtime less than to 40 while the \$lte returns less than or equal to 40.

Diving Into Read Operation

Querying Embedded Fields & Arrays

We are not limited to querying top level fields and are also able to query embedded fields and arrays. To query embedded fields and arrays is quite simple as demonstrated below:

```
$ db.movies.find( { "rating.average": { $gt: 7 } } ).pretty( )
```

We specify the path to the field that we are interested in querying the data. We must put the path within quotations marks because we use the dot (which will invalidate the syntax) to detail each embedded field within the path that leads to the field we are interested in. The above example is a single level embedded document, if we wrote e.g. rating.total.average, this is a 2 level embedded document. We can make the path as deep as we need it to be and we are not limited to one level.

We can also query arrays as seen below:

```
$ db.movies.find( { genres: "Drama" } ).pretty( )
```

```
$ db.movies.find( { genres: ["Drama"] } ).pretty( )
```

The casing is important in the query. This will return all genres that has Drama included in it. Equality in an array does not mean that Drama is the only item within the array; it means that Drama exists within the array. If we wanted an exactly only Drama within the array we would use square brackets. We can also use dot to go down embed level paths that has an array.

Diving Into Read Operation

Understanding \$in and \$nin

If there are two discrete values that we wish to check/query our data, for example runtime that is either 30 and/or 42, we can use the \$in operator. The \$in operator takes in an array which holds all the values that will be accepted to be values within the key/field.

The below example return data that have a runtime equal to 30 or 42.

```
$ db.movies.find( { runtime: { $in: [ 30, 42 ] } } ).pretty( )
```

The \$nin on the other hand is the opposite to \$in operator. It finds everything where the value is not within the set of values defined in the square brackets. The below example returns all entries where the runtime is not equal to 30 or 42.

```
$ db.movies.find( { runtime: { $nin: [ 30, 42 ] } } ).pretty( )
```

We have now explore all the Comparison operators within mongoDB and will continue to now look at logical query operators such as \$and, \$not, \$nor and \$or in the next section.

Diving Into Read Operation

Understanding \$or and \$nor

There are four different logical operators and these are \$and, \$not, \$nor and \$or operators. We would probably use the \$or logical operator more compared to the other logical operators. Below is an example of the \$or and \$nor operator in action.

```
$ db.movies.find( { $or: [ {"rating.average": { $lt: 5 } }, {"rating.average": { $gt: 9.3 } } ] } ).pretty( )  
$ db.movies.find( { $or: [ {"rating.average": { $lt: 5 } }, {"rating.average": { $gt: 9.3 } } ] } ).count( )
```

We start the filter with the \$or to tell mongoDB that we have multiple conditions and then add an array which will hold all the conditions that mongoDB will check. The or logical condition means that it will return results that match any of these conditions. We would specify our filters as we would normally would do within our find but now held within the \$or array. We can have many expressions as we want within the \$or array, in the above we have two conditions. Note: if we change pretty() for count(), this will return the total number of documents that meet the criteria rather than the document itself.

```
$ db.movies.find( { $nor: [ {"rating.average": { $lt: 5 } }, {"rating.average": { $gt: 9.3 } } ] } ).pretty( )
```

The \$nor operator is the opposite/inverse of the \$or operator. It returns where neither of the conditions are met i.e neither conditions are true the complete opposite.

Diving Into Read Operation

Understanding \$and Operator

The syntax for the \$and operator is similar to the \$or and \$nor operator. The array of documents acts as the logical conditions and will return all documents where all conditions are met. We can have as many conditions as we want. Below is an example of the \$and logical operator:

```
$ db.movies.find( { $and: [ { "rating.average": {$gt: 9} }, { genres: "Drama" } ] } ).pretty( )
```

In this example, we are trying to find all documents that are Drama with a high rating that is greater than 9. This is the old syntax and there is now a shorter syntax as seen below:

```
$ db.movies.find( { { "rating.average": {$gt: 9}, genres: "Drama" } }).pretty( )
```

The new shorter syntax does not require the \$and operator, instead we use a single document and write our conditions separating each condition with a comma. By default, MongoDB ands all key fields that we add to the filtered document. The \$and is the default concatenation for MongoDB.

The reason why we have the \$and operator is because not all drivers accepts the above syntax. Furthermore, the above shorthand syntax would return a different result when we filter on the same key elements.

If we examine the two syntax below, we would notice that they both return a different result.

```
$ db.movies.find( { $and: [ { genre:"Drama" }, { genres: "Horror" } ] } ).count( )
```

```
$ db.movies.find( { { genre: "Drama", genres: "Horror" } } ).count( )
```

When we use the second syntax, mongoDB replaces the first condition with the second and therefore it is the same as filtering for Horror genre only.

```
$ db.movies.find( { { genre: "Drama", genres: "Horror" } } ).count( )
```

```
$ db.movies.find( { genres: "Horror" } ).count( )
```

Therefore, in the scenario where we are looking for both the Drama and Horror values from the genre key element, it is recommended to use the \$and operator for mongoDB to look for both conditions to return true.

Diving Into Read Operation

Understanding \$not Operator

The \$not operator inverts the effect of a query operator. For example if we query to find movies that do not have a runtime of 60minutes as seen in the below syntax.

```
$ db.movies.find( { runtime: { $not: { $eq: 60 } } } ).count( )
```

The \$not operator is less likely to be use as it can be achieved using much simpler alternatives for example we can use the not equal operator or \$nor operator:

```
$ db.movies.find( { runtime: { $ne: 60 } } ).count( )
```

```
$ db.movies.find( { $nor [ { runtime: { $eq: 60 } } ] } ).count( )
```

There are a lot of ways for querying the inverse, however, where we cannot just simply inverse the query in another way, we have the \$not which we can use to look for the opposite.

We have now examined all four of the logical operators available within mongoDB that we can use as filters for our Read operations.

Diving Into Read Operation

Element Operators

There are two types of element operators which are `$exists` and `$type`. This allows us to query by elements within our database collection.

```
$ db.users.find( { age: { $exists: true } } ).pretty( )
```

This will check within our database and return all results where the document contains an age element/field. Alternatively, we could have queried `$exists` to be false in order to retrieve all documents that do not have age as an element/field.

We can query the `$exists` operator with other operators. In the below example we are filtering by the age element to exist and age is greater than or equal to 30:

```
$ db.users.find( { age: { $exists: true, $gte: 30 } } ).pretty( )
```

To search for a field that exists but also has a value in the field, we would query as seen below:

```
$ db.users.find( { age: { $exists: true, $ne: null } } ).pretty( )
```

The `$type` operator on the other hand, as the name would suggests, returns the document that have

The specified field element of the specified data type.

```
$ db.users.find( { phone: { $type: "number" } } ).pretty( )
```

The example above returns documents where the phone field element has values of the data type number. Number is an alias that basically sums up floats and integers. If we searched for the type of double this would also return back a document even if there are no decimal places. Since the shell is based on JavaScript, by default, a number inserted into the database will be stored as a floating point number/double because JavaScript which drives the shell does not know the difference between integers and doubles as it only knows doubles. The shell takes the number and stores it as a double even though if we have no decimal places. This is the reason why we could also search by the type of double and retrieve the documents as well.

We can also specify multiple types by passing an array. The below will look for both data types of a double and a string and return documents that match the filter condition:

```
$ db.users.find( { phone: { $type: [ "double", "string" ] } } ).pretty( )
```

We can use the type operator to ensure that we only work with the right type of data when returning some documents.

Diving Into Read Operation

Understanding Evaluation Operators - \$regex

The \$regex operator allows us to search for text. This type of query is not super performant. Regex stands for regular expression and it allows us to search for certain text based on certain patterns. Regular expressions is a huge complex topic on its own and is something not covered deeply within this MongoDB guide. Below is an example of using a simple \$regex operator.

```
$ db.movies.find( { summary: { $regex: /musical/ } } ).pretty( )
```

In this example the query will look at all the summary key field values to find the word musical contained in the value and return all matching results.

Regex is very useful for searching for text based on a pattern, however, it is not the most efficient/performant way of searching/retrieving data (the text index may be a better option and we will explore this in later chapters).

Diving Into Read Operation

Understanding Evaluation Operators - \$expr

The \$expr operator is useful if we want to compare two fields inside of one document and then find all documents where this comparison returns a certain result. Below is an example code:

```
$ use financialData
```

```
$ db.sales.insertMany( [ { volume: 100, target: 120 }, { volume: 89, target: 80 }, { volume: 200, target: 177 } ] )
```

```
$ db.sales.find( { $expr: { $gt: [ "$volume", "$target" ] } } ).pretty( )
```

In the above \$expr (expression) we are retrieving all documents where the volume is above the target. This is the most typical use case where we would use the expression operator to query the data in such a manner.

The \$expr operator takes in a document describing the expression. We can use comparison operators like gt, lt and so on – more valid expressions and which operators we can use can be found in the official documentation. We reference fields in the array rather than the number, and these must be wrapped in quotation marks along with a dollar sign at the beginning. This will tell MongoDB to look in the field and use the value in the expression. This should return two documents that meet the expression criteria.

Below is another more complex expression example:

```
$ db.sales.find( { $expr: { $gt: [ { $cond: { if: { $gte: [ "$volume", 190 ] }, then: { $subtract: [ "$volume", 10 ] }, else: "$volume" } } }, "$target" ] } } ).pretty( )
```

Not only are we comparing whether volume is greater than target but also where volume is above 190, the difference between volume and target must be at least 10. To achieve this we have to change the expression inside our \$gt operator.

The first value will be a document where we use a special \$cond operator for condition. The \$cond works in tandem with the \$expr operator. We are using an if: and then: to calculate the value dynamically. The if is another comparative operator. We are \$subtracting 10 from the volume value for all the items that are greater than or equal to 190. We use an else: case to define cases that do not match the above criteria, and in this case we would just use the volume value.

We would finally compare the value with the target to check whether the value is still greater than or equal to the target. This should return 2 documents.

As we can see from the example above, this is a very powerful command within our tool belt when querying data from a MongoDB database.

Diving Into Read Operation

Diving Deeper into Querying Arrays

There are multiple things we can perform when querying arrays and there are special operators that help us with querying arrays. If we want to search for example all documents that have an embedded sports document, we cannot use the normal queries that we have previously used for example, if we had embedded documents for hobbies that had title and frequency:

```
$ db.users.find( { hobbies: "Sports" } ).pretty( )
```

```
$ db.users.find( { hobbies: { title: "Sports" } } ).pretty( )
```

Both of these will not return any results if there are multiple fields within an embedded document.

```
$ db.users.find( { hobbies: { title: "Sports", frequency: 2 } } ).pretty( )
```

This will find any documents that meet both the criteria, however, what if we only want to retrieve all documents that have an embedded Sports document in an array, regardless of the frequency?

```
$ db.users.find( { "hobbies.title": "Sports" } ).pretty( )
```

We search for a path using dot notation. This must be wrapped in quotation. MongoDB will go through all of the hobbies elements and within each element it will dig into the document and compare title to our query value of Sports. Therefore, this will retrieve the relevant documents even if within an embedded array and there are multiple array documents. This is how we would query array data using the dot notation.

Diving Into Read Operation

Using Array Query Selector - \$size, \$all & \$elemMatch

There are three dedicated query selectors for Arrays which are \$size, \$all and \$elemMatch operators. We will examine each of these selectors and their applications.

The \$size selector operator allows us to select or retrieve documents where the array is of a certain size, for example we wanted to return all documents that had an embedded array size of 3 documents. For example:

```
$ db.users.find( { hobbies: { $size: 3 } } ).pretty( )
```

This will return all documents within the users collection where the hobbies array size is 3 documents. Note: the \$size operator takes an exact match of a number value and cannot be something like greater or less than 3. This is something MongoDB does not support yet and we will have to retrieve using a different method.

The \$all selector operator allows us to retrieve documents from an array based on the exact values without worrying about the order of the items within the array. For example if we had a movie collection where we wanted to retrieve those with a genre of thriller and action but without caring

for the order of the values, this is where the \$all array selector will help us.

```
$ db.movies.find( { genre: ["action", "thriller"] } ).pretty( )
```

```
$ db.movies.find( { genre: { $all: ["action", "thriller"] } } ).pretty( )
```

The second syntax will ensure both array elements of action and thriller exists within the genre field and ignores the ordering of these elements (i.e. ordering does not matter) whereas, the first syntax would take the order of the elements into consideration (i.e. the ordering matters).

Finally, the \$elemMatch array selector allows us to retrieve documents where one and the same element should match our conditions.

In the syntax below, this will find all documents where the hobbies has at least one document with the title of Sports and a document with a frequency greater than or equal to 3 and it does not have to be the same document/element. This would mean that a user who has the title of Sports but a frequency below 3 and a title of a different hobby but a frequency greater then 3, will match the criteria as it has at least one of each document criteria.

```
$ db.users( { $and: [ { "hobbies.title": "Sports" }, { "hobbies.frequency": { $gte: 3 } } ] } ).pretty( )
```

To ensure we retrieve all documents where the hobbies is Sports and its frequency is greater than or equal to 3 is returned, the \$elemMatch operator is useful:

```
$ db.movies.find( { hobbies: { $elemMatch: { title: "Sports", frequency: { $gte: 3 } } } } ).pretty( )
```

Understanding Cursors

The `find()` method unlike the `findOne()` method yields us with a cursor object. Why is this cursor object important? If our client communicates with the MongoDB server and we potentially retrieve thousands if not millions of documents with our `find()` query depending on the scale of the application. Retrieving all these results is very insufficient because all these documents have to be fetched from the database, they have to be sent over the wire and then loaded into memory in the client application. These are the three things that are not optimal. In most cases we would not need all the data at the same time; therefore, `find` gives us a cursor.

A cursor is basically a pointer which has the query we wrote stored and can go back to the database as request the next batch. We therefore work with batches of data and we fetch documents over the wire one document at a time. The shell by default takes the cursor and retrieves the first 20 documents, it then fetches in batches of 20 documents.

If we write our own application with a MongoDB driver, we have to control that cursor manually to make sure that we get back our results. The cursor approach is beneficial because it saves on resources and we only load a batch of documents rather than all the documents from a query.

Diving Into Read Operation

Applying Cursors

When using the find command in the shell this will display the first 20 documents. We can use the “it” command in shell to retrieve the next 20 batches of documents and keep using this command until we have exhausted the cursor i.e. there are no more documents to load. The “it” command will not work with the MongoDB drivers. Instead most drivers will have a next() method we can call instead. If we use the next() method within the shell, this will retrieve only one document and there will be no “it” command to retrieve the next document. If we run the command again this will restart the find query retrieving the first document again.

```
$ db.movies.find( ).next( )
```

Since the shell uses JavaScript, we can use JavaScript syntax to store the cursor value in a variable. We can then use the next() method to cycle through the next document on the cursor which will retrieve the next document continuing on from the last cursor value.

```
$ const dataCursor = db.movies.find( )
```

```
$ dataCursor.next( )
```

There are other cursor methods available to us in MongoDB that we can use on our find() query.


```
$ dataCursor
```

```
$ it
```

This will return the first 20 batch of documents. Using the `it` command will retrieve the next 20 documents i.e. the default shell behaviour for cursors.

```
$ dataCursor.forEach( doc => { printjson(doc) } )
```

The `forEach()` method will vary on the driver we are using, but in JavaScript the `forEach()` method takes in a function that will be executed for every element that can be loaded through the cursor. In javascript we get a document which is provided by the `forEach` method which is passed in as the input and then our arrow function provides the body what we want to do. In the above example we are using the `printjson()` method which is provided by the shell to output the document.

This will cycle through all the remaining documents inside of the cursor (this will exclude any documents we have already searched for i.e. using the `next()` method or `find()` method). The `forEach` will retrieve all the remaining documents and there will be no “`it`” command to fetch the next documents as we would have exhausted all the documents in the cursor.

```
$ dataCursor.hasNext( )
```

The `hasNext()` method will return true or false to indicate if we have/have not exhausted the cursor. We can create a new variable to reset the cursor for `const` variables or if we used `let` or `var` variables we can re-instantiate the original variable again to reset the cursor to the beginning.

We can learn more on cursors and the shell or the MongoDB drivers on the official MongoDB documentation:

<https://docs.mongodb.com/manual/tutorial/iterate-a-cursor/>

Diving Into Read Operation

Sorting Cursor Results

A common operation is to sort the data that we retrieve. We are able to sort by anything whether it is a string sorted alphabetically or a number sorted by numeric value. To sort the data in MongoDB we would use the `sort()` method on our `find` function. Below is an example:

```
$ db.movies.find( ).sort( { "rating.average": 1 } ).pretty( )
```

The `sort` takes in a document to describe how to sort the retrieved data. We can sort by a top level document field or an embedded document field. The values we use to sort describe the direction to sort the data i.e. `1` means ascending (lowest value first) and `-1` means descending (highest value first). We are not limited to 1 sorting criteria for example we want to sort by the average ratings first but then we want to sort by the runtime:

```
$ db.movies.find( ).sort( { "rating.average": 1, runtime: -1 } ).pretty( )
```

Diving Into Read Operation

Skipping & Limiting Cursor Results

We are able to skip a certain amount of elements. Why would we want to skip elements? If on our application or web app we implement pagination where users can view results distributed across multiple pages (e.g. 10 elements per page) because we do not want to show all results on one page. If the user switches to a page e.g. page 3, we would want to skip the first 20 results to display the results for page 3.

The skip method allows us to skip cursor results. Below is an example of skipping by 20 results. Skipping allow us to move through our dataset.

```
$ db.movies.find( ).sort( { "rating.average": 1, runtime: -1 } ).skip(20).pretty( )
```

The limit method allows us to limit the amount of elements the cursor should retrieve at a time and then also means the amount of documents we can then move through a cursor. Limit allows us to retrieve a certain amount of documents but only the amount of documents we specify.

```
$ db.movies.find( ).sort( { "rating.average": 1, runtime: -1 } ).skip(100).limit(10).pretty( )
```

We can have the sort, skip and limit methods in any order we want and this will not affect the sorting as mongoDB will automatically do the actions in the correct order.

Diving Into Read Operation

Using Projections to Shape Results

How can we shape the data which we get back from our find into the form we need it? When we use the find function this retrieves all the data from the retrieved document. This is not only too much redundant data transferred over the wire but also makes it hard to work with the data if we have to manually parse all the data. Projection allows us to control which data is returned from our Read Operation.

Below is an example code for projecting only the name, genre, runtime and rating from the returned results and all other data on the document does not matter to us.

```
$ db.movies.find( { }, { name: 1, genre: 1, runtime: 1, rating: 1 } ).pretty( )
```

In order to perform a projection we need to pass a second argument to the find function. If we do not want to specify any filter criteria for the first criteria of find, we would simply add an empty document as seen above. The second argument allows us to configure how values are projected i.e how we extract the data fields. We name the field and pass the value of 1. All fields that we do not mention with a 1 or we explicitly add with a 0 will be excluded by default. This will retrieve a reduced version of the document, but the id will always be included in the results even if we do not specify it within the projection. If we want to exclude the id we must explicitly exclude it as seen

in the below example:

```
$ db.movies.find( { }, { name: 1, genre: 1, runtime: 1, rating: 1, _id: 0 } ).pretty( )
```

This will retrieve the data in the projection and explicitly exclude the `_id` value from the results. This is only required for the `_id` field and all other fields are implicitly set to 0 to exclude by default on projections.

We are also able to project on embedded documents for example we are only interested in the time from the schedule embedded document and not the day field. We would use the path notation to select the embedded document to project as seen below:

```
$ db.movies.find( { }, { name: 1, genre: 1, runtime: 1, "schedule.time": 1 } ).pretty( )
```

Projections can also work with arrays and help with array data that we include in our returned find query results. For example, if we want to project Drama only from the genres, we would first filter the data by the criteria of all documents with an array containing Drama in the Genres and then use the projection argument to display only the Drama from the array:

```
$ db.movies.find( { genres: "Drama" }, { "genres.$": 1 } ).pretty( )
```

The special syntax of `$` after the genres will tell MongoDB to project on the one genre it found.

Now the document may have had more than Drama within the array in the retrieved document, but we have told MongoDB to only fetch the Drama and only output that result because that is the only data we are interested in retrieving. The items behind the scenes may have more data, just as they have other fields too. However, with the \$ syntax will find the first matching document from our criteria query which in the above scenario was Drama.

If we had a more complex criteria whereby we are searching for all documents with both Drama and Horror, the \$ projection syntax will return Horror because the Horror is the first matching criteria in the below example. The \$all requires a match when Drama is present but the matching is when Horror is also present and therefore the projection will display Horror as it is technically the first matching criteria because Drama didn't match anything.

```
$ db.movies.find( { genres: { $all: [ "Drama", "Horror" ] } }, { "genres.$": 1 } ).pretty( )
```

There may be cases where we want to project items from an array in our document that are not items we queried for. In the below example we query to retrieve all documents with the value of Drama, but we want to project Horror only. Using the \$elemMatch operator allows us to do this:

```
$ db.movies.find( { genres: "Drama" }, { genre: { $elemMatch: { $eq: "Horror" } } } ).pretty( )
```

The filter criteria and projection can be totally unrelated as seen in the below example:

```
$ db.movies.find( { "rating.average": { $gt: 9 } }, { genre: { $elemMatch: { $eq: "Horror" } } } ).pretty( )
```

Diving Into Read Operation

Understanding \$slice

The final special projection relating to arrays is called the \$slice operator. This operator returns the first x amount of items from the array. In the below syntax example we are projecting the first 2 array items within the genres field.

```
$ db.movies.find( { "rating.average": { $gt: 9 } }, { genres: { $slice: 2 }, name: 1 } ).pretty( )
```

The documents may have more genres assigned to them but we only see the first two items in the array because we used the \$slice: 2 to return only 2 items. The number can be any value to return any number of items from the array. The alternative method of slice is to use an array form.

```
$ db.movies.find( { "rating.average": { $gt: 9 } }, { genres: { $slice: [ 1, 2 ] }, name: 1 } ).pretty( )
```

The first element in the slice array is the amount of elements in the array which we skip e.g. we skip one item. The second element in the slice array is the amount of data we want to limit it to e.g. we want to limit it to two items. This will return item 2 and item 3 from the array and skip item 1 when it projects the results (i.e. we projected the next two items in the array).

We have many ways of controlling what we see using the 1 and 0 for normal fields and using the \$, \$elemMatch and \$slice to control which element in the array are projected in our results.

We have now completely explored how to fully control what we fetch with our filtering criteria and then control which fields of the found document to include in our displayed result sets using projections.

Useful Links:

<https://docs.mongodb.com/manual/reference/method/db.collection.find/>

<https://docs.mongodb.com/manual/tutorial/iterate-a-cursor/>

<https://docs.mongodb.com/manual/reference/operator/query/>

Diving Into Update Operation

Updating Fields with `updateOne()`, `updateMany()` and `$set`

An update operation requires two pieces of information (two properties):

1. Identify the document that should be updated/changed (i.e. the filter)
2. Describe how should it be updated/changed

For identifying the document we can use any of the filters we could use for finding documents and do not necessarily need to use the `_id` value. Using the `_id` is a common use for updating documents as it will guarantee the correct document is being updated.

The `updateOne()` method simply takes the first document that matches the filter criteria and updates it even if multiple documents may match the criteria. The `updateMany()` method on the other hand will take the criteria/filter and update all documents that match.

```
$ db.users.updateOne( { _id: ObjectId("5b9f707ead7") }, { $set: { hobbies: [ { title: "Sports", frequency: 5 }, { "Cooking", frequency: 3 }, { title: "Hiking", frequency: 1 } ] } } )
```

```
$ db.users.updateMany( { "hobbies.title": "Sports" }, { $set: { isSporty: true } } )
```


The second argument/parameter is how to update the document and the various update operators can be found on the official MongoDB documentation:

<https://docs.mongodb.com/manual/reference/operator/update/>

The \$set takes in a document where we describe some fields that should be changed or added to the existing document. The updateOne example overwrites the existing hobbies array elements with the new hobbies array elements. The console will provide a feedback when it updates the document:

```
{ "acknowledged" : true, "matchedCount" : 1, "modifiedCount" : 1 }
```

If we ran the exact updateOne command again the console will still show a matchedCount of 1 but the modifiedCount will be 0 as no document data would have been modified because it is exactly the same as the existing value.

```
{ "acknowledged" : true, "matchedCount" : 1, "modifiedCount" : 0 }
```

If we were to find all the documents using the find() command, we would notice the user would still have the same _id and other document fields. This is because the \$set operator does not override the existing values, instead it simply defines some changes that MongoDB evaluates and then if they make sense it changes the existing document by adding or overwriting the second argument fields. All the existing fields are left untouched. The \$set operator by default simply adds or edits the fields specified in the update command.

Diving Into Update Operation

Updating Multiple Fields with \$set

In the previous section we demonstrated the \$set operator used to update one field at a time in a document. It is important to note that the \$set is not limited to updating one field in a document but can update multiple fields within a document as seen in the below example whereby we add a field of age and another field of phone:

```
$ db.users.updateOne( { _id: ObjectId("5b9f707ead7") }, { $set: { age: 40, phone: 07933225211 } } )
```

The console again will output when a document has matched the filter and have been modified whether the update modification was a overwrite or adding new fields to the matched document.

```
{ "acknowledged" : true, "matchedCount" : 1, "modifiedCount" : 1 }
```

The \$set operator can add fields, arrays and embedded documents inside and outside of an array.

Diving Into Update Operation

Incrementing & Decrementing Values

The update operator allows us not only allows us to simply setting some values, but it can also

increment or decrement a number for us. For example, if we wanted to update a users age without using the \$set operator as we would not necessarily know ever users age, rather we would want MongoDB to perform a simple common transformation of taking the current age and then recalculate the new age and then issue the update request.

MongoDB has a built in operator to allow us to perform common increment and decrement operations using the \$inc operators.

```
$ db.users.updateOne( { name: "Emanuel" }, { $inc: { age: 1 } } )
```

This will increment the existing age field value by one. Note we could choose a different number such as 5 and increment the number by 5. To decrement a filed value we would continue to use the \$inc operator but use a negative number instead.

```
$ db.users.updateOne( { name: "Emanuel" }, { $inc: { age: -1 } } )
```

Note we can perform multiple different things in the same update such as increment certain fields while setting new fields/editing existing fields, all within the second update parameter.

```
$ db.users.updateOne( { name: "Emanuel" }, { $inc: { age: 1 }, $set: { gender: "M" } } )
```

If we tried to increment/decrement a value as well set the same field value, this will give us an error:

```
$ db.users.updateOne( { name: "Emanuel" }, { $inc: { age: 1 }, $set: { age: 30 } } )
```

This would error in the console because updating the path will cause a conflict because we have two update operations working on the same field and this is not allowed in MongoDB and will fail providing a message in the shell something like the below:

```
$ WriteError: Updating the path 'age' would create a conflict at 'age': ...
```

Diving Into Update Operation

Using \$min, \$max and \$mul Operators

The \$min, \$max and \$mul are other useful operators available to us. Below are scenarios and the syntax to overcome various update problems.

Scenario 1: We want to set the age to 35 only if the existing age is higher than 35. We would use the \$min operator because this will only change the current value if the new value is lower than existing value. Note: this will not throw any errors if the existing value is higher than the new value, it will simply not update the document.

```
$ db.users.updateOne( { name: "Chris" }, { $min: { age: 35 } } )
```

Scenario 2: We want to set the age to 38 only if the existing age is lower than 38. We would use the \$max operator which is the opposite of the \$min operator. Again this will not throw any error for existing values is lower than the update value as it will simply ignore the update.

```
$ db.users.updateOne( { name: "Chris"}, { $max: { age: 38 } } )
```

Important Note: the modifiedCount in the terminal will show as 0 to show if no update occurred.

Scenario 3: We want to multiply the existing value with a multiplier. We would use the \$mul operator which stands for multiply to perform this type of update operation.

```
$ db.users.updateOne( { name: "Chris"}, { $mull: { score: 1.1 } } )
```

This will multiply the existing score value by the multiplier of 1.1 to update the score document with the new value.

Diving Into Update Operation

Getting Rid of Fields

If we want to update records to drop a field based on a certain criteria(s). There are two solutions to this problem and below is an example syntax to drop all phone numbers for users that are isSporty.

```
$ db.users.updateMany( { isSporty: true }, { $set: { phone: null } } )
```

We could use the \$set operator to set the phone to null. This will not drop the field but it would mean that the field has no value which we can use in our application to not display the phone data.

The alternative solution is to use the special operator of \$unset to truly get rid of a field from a document.

```
$ db.users.updateMany( { isSporty: true }, { $unset: { phone: "" } } )
```

The value we use with phone (or key field) is totally up to us but typically set to "" which represents empty. The key would be ignored in the update as the important part of the \$unset operator document is the field name we wish to drop.

Diving Into Update Operation

Renaming Fields

Just as we have the \$unset operator to drop a field from a document, we are also able to rename fields using the \$rename operator. We can leave the first update argument empty to target all documents and update the field name. The \$rename document takes in the field name we want to rename and the key as a string of the new field name value. This will only update all documents that has a field called age to the new updated field name.

```
$ db.users.updateMany( { }, { $rename: { age: "totalAge" } } )
```

Diving Into Update Operation

Understanding The Upsert Option.

If we wanted to update some documents but we were uncertain if the document existed or not. For example we have an application but we did not know if the data was saved in the database yet and if it wasn't saved yet, we now want to create a new document and if it did exist we want to overwrite/update the existing document.

In this scenario if Maria did not exist as a document, nothing will happen other than a message in the terminal to tell us no document was updated. Instead, we would want MongoDB to create the document for us instead of having us manually doing this.

```
$ db.users.updateOne( { name: "Maria" }, { $set: { age: 29, hobbies: [ { title: "Cooking", frequency: 3 } ], isSporty: true } } )
```

To allow MongoDB to update or create documents for us, we would pass a third argument option called `upsert` and set this to `true` (by default this is set to `false`). `Users` is a combination of `update` and `insert` and will work with both `updateOne` and `updateMany` methods. If Maria does not exist it will create a new document and it will also include the name: "Maria" field automatically for us.

```
$ db.users.updateOne( { name: "Maria" }, { $set: { age: 29, hobbies: [ { title: "Cooking", frequency: 3 } ], isSporty: true } }, { upsert: true } )
```


Diving Into Update Operation

Understanding Matched Array Elements.

Example scenario: we want to update all users document where the person has a hobby of sports and the frequency is greater or equal to three. The hobby field has an array of embedded documents.

```
$ db.users.find( { $and: [ { "hobbies.title": "Sports", { "hobbies.frequency": { $gte: 3 } } } ] } ).pretty( )
```

This syntax will find all users which has hobbies title of Sports and hobbies that as a frequency that is greater or equal to three. This will find documents even where the Sports frequency is below three so long as there is another embedded hobbies document which has a frequency greater or equal to three.

The correct query is to use the \$elemMatch operator which will look at the same embedded document for both criteria:

```
$ db.users.find( { hobbies: { $elemMatch: { title: "Sports", frequency: { $gte: 3 } } } } ).pretty( )
```

Now that we know the correct query to find the documents we wish to update, the question now becomes how do we update that embedded array document only. So essentially we know which

overarching document we want to change but we want to change something exactly within that document found in the array.

The \$set operator, by default will apply the changes to the overall document and not the document in the array we found. We would use the \$set operator and then select the array field and use .\$ as the syntax. This will automatically refer to the element matched in our filter criteria (first argument to the update command). We can define the new value after the colons.

```
$ db.users.updateMany( { hobbies: { $elemMatch: { title: "Sports", frequency: { $gte: 3 } } } }, { $set: { "hobbies.$": { title: "Sports", frequency: 3 } } } )
```

Note this will update all matching documents to the updated changes. However, if we only want to add a new field to all matching documents the syntax would be to add a dot after the .\$ as seen below:

```
$ db.users.updateMany( { hobbies: { $elemMatch: { title: "Sports", frequency: { $gte: 3 } } } }, { $set: { "hobbies.$.highFrequency": true } } )
```

The above syntax will find all documents which match the embedded document criteria and update that embedded document to add a new field/value of highFrequency: true if the highFrequency field did not exist (if it did, it would simply update the existing field value). The \$set works exactly as it did before the only difference is adding the special placeholder within the array path to quickly get access to the matched array element.

Diving Into Update Operation

Updating All Array Elements.

Example Scenario: The below find method returns the overall person document where the filter matched and does not return the document we filtered on only, the filter is just a key factor for returning to us the overall document.

```
$ db.users.find( { "hobbies.frequency": { $gt: 2 } } ) .pretty( )
```

This will also retrieve a document with a frequency lower than two provided at least one embedded document has a frequency above two to match the filter. Now that we have found all documents meeting the criteria but not all array documents fulfilled our filter criteria. However, we want to change all embedded documents in the hobbies array that did fulfil the filter criteria only.

```
$ db.users.updateMany( { "hobbies.frequency": { $gt: 2 } }, { $set: { "hobbies.$.goodFrequency: true } } )
```

Again we can use the \$set operator, but we want to change all the matched hobbies with a frequency above two. The "hobbies.\$" syntax we saw in the previous section only edits hobby for each person and if there is multiple matching hobbies per person, it will not update them all but only the first element within the array.

Now in order to update all elements in an array, we would use a special placeholder in MongoDB that is the `.$[]` which simply means update all elements. We can use the dot notation after the `.$[]` to select a specific field in the array document.

```
$ db.users.updateMany( { totalAge: { $gt: 30 } }, { $inc: { "hobbies.$[].frequency": -1 } } )
```

This will update all users that has a totalAge greater than 30 and decrement the hobbies frequency by 1. The `.$[]` syntax is used to update all arrays elements.

Diving Into Update Operation

Finding and Updating Specific Fields

Continuing on from the last section, we were able to use the `.$[]` notation to update all elements within the array per person. We can now build on this notation to update specific fields and below is the solution to the previous problem.

```
$ db.users.updateMany( "hobbies.frequency": { $gt: 2 } }, { $set: { "hobbies.$[el].goodFrequency": true } }, { arrayFilters: [ { "el.frequency": { $gt: 2 } } ] } )
```

Note: Within the `el` within the square bracket is an identifier which we could have named as anything. For the third argument in our update method we would use the `arrayFilters` option to define the identifier/filter elements. The identifier does not need to be the same as the filter criteria for example we could filter by age but use the `$set` to identifier to update based on a

complete different filter such as the frequency greater than two. The `arrayFilter` can have multiple documents for each identifier.

This would now update specific array elements that meet the identifier `arrayFilter` criteria. Note that the third argument of `updateOne` and `updateMany` is an options argument whereby we previously used `upsert` as an option to update/insert documents. We can also use `arrayFilters` to provide a filter criteria for our identifiers.

Note: if an identifier is used we must use `arrayFilter` option to define the identifier filter criteria otherwise the update method will fail as `mongoDB` will not know what to do with the identifier.

Diving Into Update Operation

Adding Elements to Arrays

If we want to add elements onto an array for a document instead of using `$set` operator (which we can still use to update fields), we can use `$push` to push a new element onto the array. The `$push` operator takes in a document where we describe firstly the array we wish to push to and then the element we want to push/add to the existing array documents.

```
$ db.users.updateOne( { name: "Maria" }, { $push: { hobbies: { title: "Sports", frequency: 2 } } } )
```

The \$push operator can be used with more than one document to be added. We use the \$each operator which takes in an array of multiple documents that should be added/pushed.

```
$ db.users.updateOne( { name: "Maria" }, { $set: { $push: { hobbies: { $each: [ { title: "Running", frequency: 1 }, { title: "Hiking", frequency: 2 } ] } } } } )
```

There are two options sibling options we can add to the above \$each syntax. The first is the \$sort operator. Technically, we are still in the same object where we have the \$each operator i.e. it is a sibling to \$each. The sort describes how the elements in the array should be sorted before they are pushed into hobbies.

```
$ db.users.updateOne( { name: "Maria" }, { $set: { $push: { hobbies: { $each: [ { title: "Running", frequency: 1 }, { title: "Hiking", frequency: 2 } ], $sort: { frequency: -1 } } } } } )
```

This will sort the hobbies array by frequency in a descending order i.e. having the highest frequency first.

The second sibling is the \$slice operator which allows us to add only a certain number of element. We can use this in conjunction with the \$sort operator as seen below.

```
$ db.users.updateOne( { name: "Maria" }, { $set: { $push: { hobbies: { $each: [ { title: "Running", frequency: 1 }, { title: "Hiking", frequency: 2 } ], $sort: { frequency: -1 }, $slice: 1 } } } } )
```

In this example the slice is taking only the first element after the sort to push to the hobbies array. The sort is on the overall array i.e. new and existing elements and not just the elements we add.

Diving Into Update Operation

Removing Elements from Arrays

Not only are we able to push elements to an array but we are also able to pull elements from an array using the \$pull operator, as demonstrated below.

```
$ db.users.updateOne( { name: "Maria" }, { $pull: { hobbies: { title: "Hiking" } } } )
```

The \$pull operator takes in a document where we describe what we want to pull from the array. In the above we are pulling from the hobbies array based on a condition i.e. pull every element where the title is equal to Hiking. We do not need to only use equality conditions but we can also use all the normal filter operators that we have seen before such as the greater than or less than operator.

Sometimes we may wish to remove the last element from an array and have no specific filter criteria. We would use the \$pop operator and use a document to define the name of the field of which we want to pop. The -1 defines to pop the first element and the 1 defines to pop the last element from the array.

```
$ db.users.updateOne( { name: "Chris" }, { $pop: hobbies: 1 } )
```

Diving Into Update Operation

Understanding the \$addToSet Operator

The final update command we will explore is the \$addToSet operator.

```
$ db.users.updateOne( { name: "Maria" }, { $addToSet: { hobbies: { title: "Hiking", frequency: 2 } } } )
```

The difference between \$addToSet and \$push operator, the \$push operator allows us to push duplicate values whereas the \$addToSet does not allow for this. It is important to note the console will not error but simply show that no document was updated with the change. Always remember that \$addToSet operator adds unique values only.

This concludes the Update Operations available to us in MongoDB. We now understand the three arguments we can pass to both the updateOne and updateMany commands which are:

1. Specify a filter (query selector) using the same operators we know from find() command.
2. Describe the updates via \$set or other update operators.
3. Additional Options e.g. \$upsert or \$arrayFilters to the update operation.

In the official documentation we can view all the various operations available to us:

<https://docs.mongodb.com/manual/tutorial/update-documents/>

Diving Into Delete Operation

Understanding deleteOne() and deleteMany()

To delete a single document from a collection we would use the deleteOne() command. We need to specify a query selector/filter. The filter we specify here is exactly the same as we would use for the the finding and updating documents. We simply need to narrow down the document we want to delete. DeleteOne will delete the first document that matches the criteria.

```
$ db.users.deleteOne( { name: "Chris" } )
```

We can use the deleteMany() command to delete all documents where the query selector/filter criteria has been met. Below are two examples.

```
$ db.users.deleteMany( { totalAge: {$gt: 30}, isSporty: true } )
```

```
$ db.users.deleteMany( { totalAge: {$exists: false}, isSporty: true } )
```

Note: we can add as many query selectors as we want to narrow down the document(s) we wish to delete from the database.

Diving Into Delete Operation

Deleting All Entries in a Collection

There are two approaches to deleting all entries in a collection. The first method is to reach out to the collection and execute the `.deleteMany()` command and pass an empty document as the argument. This argument is a filter that matches every document in the collection and therefore will delete all entries within the collection.

```
$ db.users.deleteMany( { } )
```

The alternative approach is to delete the entire collection using the `.drop()` command on the specified collection. This will return true if successfully dropped a collection.

```
$ db.users.drop( )
```

When creating an application it is very unlikely we would drop collections. Adding and dropping collections is more of a system admin task. We can also drop an entire database using the `dropDatabase` command. We would then use the `use` command followed by the database to navigate to the desired database collection.

```
$ db.dropDatabase( )
```

<https://docs.mongodb.com/manual/tutorial/remove-documents/>

Working with Indexes

What are Indexes and why do we use them?

An index can speed up our find, update or delete queries i.e. all the queries where we are looking for certain documents that should match some criteria.

```
$ db.products.find( { seller: "Marlene" } )
```

If we take a look at this find query, we have a collection of documents called products and we are searching for a seller called Abbey. Now by default if we don't have an index on the seller set, MongoDB will go ahead and do a so-called collection scan. This simply means that MongoDB to fulfil this query will go through the entire collection, look at every single document and see if the seller equals "Marlene" (equality). As we can imagine, for a very large collection with thousands or millions of document, this can take a while to complete. This is the default or only approach MongoDB can take when there are no indexes setup in order to retrieve maybe two documents out of the thousands of documents in the collection.

We can create an index and an index is not a replacement for a collection but rather an addition. We would create an index for the seller key of the product collection and that index then exists additionally to the collection and the index is essentially an ordered list of all the values that are stored in the seller key for all the documents.

It is not an ordered list of the documents, it is just the values for the field for which we created that index. Also it is not just an ordered list of the values, every value/item in the index has a pointer to the full document it belongs to.

This allows MongoDB to perform a so-called index scan to fulfil this query. This means MongoDB will see that for seller, such an index exists and it therefore simply goes to that seller index and can quickly jump to the right values because unlike for the collection, it knows the values are sorted by that key. This means if we are searching for a seller starting with M, it does not need to search through the first few records. This allows MongoDB very efficiently go through that index and find the matching product because of the ordering and the pointer that every items within the index has. So MongoDB finds the value for the query and then finds the related document to return.

This is how an index works in MongoDB and also answers why we would use indexes because creating indexes drastically speeds up our queries. However, we also should not overdo it with indexes. Lets take the example of a Products collection which has a `_id`, name, age and hobbies fields. We could create a index for all four fields and we would have the best performance because no matter what we look for, we have an index and can query for every field efficiently which will speed our find queries. Having said this, index does not come without cost. We would have to pay some performance cost on inserts because that extra index that has to be maintained would need to be updated with every inserts. This is because we have an ordered list of elements with pointers

to the documents. So if we add a new document, we also have to add a new element to the index. This may sound simple and it would not take super long, but if we have 10 indexes for our document in our collection, we would have to update all 10 indexes for every insert. We may then quickly run into some issues because we will have to do a lot of work for all these fields for every insert and for every update too. Therefore, indexes do not come for free and we have to figure out which indexes makes sense to have and which indexes don't.

We are now going to explore indexes in more detail and look at all the type of indexes that exist in MongoDB and how to measure whether an index makes sense or does not make sense to have.

Working with Indexes

Adding a Single Field Index?

To determine whether an index can help us in our find query, MongoDB provides us with a nice tool that we can use to analyse how it executed the query. We can chain the explain method to our normal query. This method works with find, update and delete commands but not for inserts (i.e. it works for methods that narrow down documents).

```
$ db.contacts.explain( ).find( { "dob.age": { $gt: 60 } } )
```

This will provide a detailed description of what MongoDB did to derive to the results.

MongoDB thinks in the so-called plans and plans are alternatives it considers for executing the query and in the end it will find the winning plan. The winning plan is simply what mongoDB did to get to our results. Without indexes, a full scan is always the only thing mongoDB can do. However, if there were alternatives and they were rejected then this will appear in the rejectedPlans array.

To get a more detailed report we can run the same command but passing in an argument:

```
$ db.contacts.explain( "executionStats" ).find( { "dob.age": { $gt: 60 } } )
```

The executionStats provides a detailed output for our query and how the results were returned. This will show things such as executionTimeMillis which is the time it took to execute the query in milliseconds and totalDocsExamined which shows the number of documents that needed to be scanned in order to return our query results. The larger the gap between the totalDocsExamined and nReturned values, this shows how inefficient the query is.

To add an index on a collection, we would use the createIndex method passing in a document. The first value (key) in the document is the name of the field we want to create a index on (this can be on top level fields as well as embedded field names). The value is whether mongoDB should create a list of values in the field in an ascending (1) or descending (-1) order.

```
$ db.contacts.createIndex( { "dob.age": 1 } )
```

Once we run the command we should see in the terminal that the index has been created.

```
{  
  "createdCollectionAutomatically" : false  
  "numIndexesBefore" : 1  
  "numIndexesAfter" : 2  
  "ok" : 1  
}
```

If we were to run the above explain command again on our collection, we should notice the executionTimeMillis for the same query has been sped up. We should also see two execution stages the first being an index scan (IXSCAN). The index scan does not return the documents but the keys/ pointers to the document. The second stage is the fetch (FETCH) which will take the pointers returned from the index scan and reach out to the actual collection and then fetch the real documents. We would notice that MongoDB would only have to look at a reduced number of documents to return the documents from our query.

This is how an index can help us to speed up our queried searches and how to use the explain method to determine whether an index should be used in our collection to speed up the query.

Working with Indexes

Indexes Behind the Scenes

What does `createIndex()` method do in detail?

Whilst we can't really see the index, we can think of the index as a simple list of values and pointers to the original document.

Something like this (for the "age" field):

(29, "address in memory/ collection a1")

(30, "address in memory/ collection a2")

(33, "address in memory/ collection a3")

The documents in the collection would be at the "addresses" a1, a2 and a3. The order does not have to match the order in the index (and most likely, it indeed won't).

The important thing is that the index items are ordered (ascending or descending - depending on how we created the index). The syntax of `createIndex({age: 1})` creates an index with ascending sorting while the syntax of `createIndex({age: -1})` creates an index with descending sorting.

MongoDB is now able to quickly find a fitting document when we filter for its age as it has a sorted list. Sorted lists are way quicker to search because we can skip entire ranges (and don't have to look at every single document).

Additionally, sorting (via `sort(...)`) will also be sped up because you already have a sorted list. Of course this is only true when sorting for the age.

Working with Indexes

Understanding Index Restriction

In the previous section we created an index which sped up our query when looking for people with an age greater than 60. However, if we run the same query but find people older than 20, we will notice that the execution time is higher than it was for people above the age of 60.

To drop an index from our collection we would use the `dropIndex` method and pass in the document that we created to create the index.

```
$ db.contacts.dropIndex( { "dob.age": 1 } )
```

If we were to run a full scan against our collection, we will notice that the query is much faster than

having an index. The reason for why the query is much faster is because we have saved a step from going through the index. If we have a query that will return a large portion or the majority of our documents, an index can actually be slower because we have an extra step to go through almost the entire index list and we then have to go to the collection and get all these documents. If we do a full scan, we do not have this extra step of going through the collection to get the documents because with a full collection scan we already have all the documents in memory and an index doesn't offer us any more because it will only be an extra step.

Important note: if we have queries that regularly return the majority or all of our documents, an index will not help us and it might even slow down the execution. This is the first important note to keep in mind (a first restriction) when planning our queries and whether or not to use indexes.

If we have a dataset where our queries typically return a fraction like 20% or lower than that of the documents, then indexes will certainly always speed up our queries. If we have a lot of queries that give us back all the documents or close to all the documents, then indexes can not do much for us. The whole point of indexes is to quickly get to a narrow subset of our document list and return the documents from that index.

Working with Indexes

Creating Compound Indexes

Not only can we have indexes on fields that have number values but we can also have indexes on fields that have text values (*both can be sorted*). We cannot create indexes for booleans as we only have two kind of values i.e. true and false and the change of index on booleans will not speed up our queries. Below is an example of creating an index on a text field.

```
$ db.contacts.createIndex( { gender: 1 } )
```

Now the above index would not make too much sense for an index because gender has two values of Male and Female and would probably return more than half the results. However, if we want to find as an example all people who are older than 30 and are male, we can create a so called compound index.

```
$ db.contacts.createIndex( { "dob.age": 1, gender: 1 } )
```

The order of the two fields in our createIndex method do matter because a compound index simply is an index with more than one field. This will store one index where each entry in the index is now not on a single value but on two combined values. This does not create two indexes and this is really important to note with compound indexes, it creates one index where every element is a connected value (it creates a pair value for example in the above this is a pair of the age and gender values).

The order of the fields defines which kind of pairs MongoDB will create in our compound index (for example does MongoDB create a 31 male index or a male 31 index – this will be important for our queries).

There are two queries we can now run which will take advantage of the compound index. The first is to find based on age and gender:

```
$ db.contacts.explain( ).find( { "dob.age": 35, gender: "male" } )
```

This will perform an index scan with our index name (*the index name is auto-generated e.g. "indexName": "dob.age_1_gender_1"*)

The second query that can utilise the compound index is a query on the age only:

```
$ db.contacts.explain( ).find( { "dob.age": 35 } )
```

This will also use the same compound index we created for the index scan even though we never specified to search for the gender. Compound indexes can be used from left to right, but the left must always be used in the search i.e. a find query on the gender alone will not work.

```
$ db.contacts.explain( ).find( { gender: "male" } )
```

The above query would use a full collection scan and not the index scan using the compound index.

The compound indexes are grouped together, the first field (left) will be ordered whereas the other fields (right) will not be ordered. We can have a compound index with more than 2 fields but up to a maximum of 31 field. However, we cannot utilise the compound index without the first field.

These are the restrictions we have on compound indexes but compound indexes allows us to speed up queries that uses multiple values.

Working with Indexes

Using Indexes for Sorting

Now that we have had a look at the basics of indexes, it is important to know that indexes are not only used for narrowing our find queries but they can also help with sorting. Now that we have a sorted list of elements of the index, MongoDB can utilise this in case we want to sort in the same way that the index list is sorted.

```
$ db.contacts.explain( ).find( { "dob.age": 35 } ).sort( { gender: 1 } )
```

In the above we can find people with the age of 35 but sort them by gender in an ascending order. We will notice that this will use an index scan for both the gender and age even though we filtered by age only. It uses the gender information for the sorting. Since we already have an ordered list of values, MongoDB can utilise this to quickly give back the order of documents we need.

It is important to understand that if we are not using indexes and we do a sort on a large amount of documents, we can actually timeout because MongoDB has a threshold of 32mb in memory for sorting. If we have no index, MongoDB will essentially fetch all our documents into memory and do the sort there and for very large collections and large amount of fetched documents, this can be too much to then sort.

Sometimes we would need indexes not only to speed up the query but also to be able to sort at all. This is not a problem for small dataset but where we fetch so many documents that an in-memory sort which is the default is just not possible and we then need an index which is already sorted so that MongoDB does not have to sort in memory but can take the order we have in the index.

Important Note: MongoDB has a threshold of 32mb which it reserves in memory for the fetched document and sorting them. This is the second important note to keep in mind as to whether or not to create an index.

Working with Indexes

Understanding the Default Index

When creating an index it would seem like that there is an index already existing in our collection. To be able to see all indexes that exists for a collection we can use the `getIndexes` command.

```
$ db.contacts.getIndexes( )
```

This command will print all the indexes we have on that collection within the shell. We will notice, if we have created new indexes on our collection that there are two indexes. The first index on the `_id` field is a default index mongoDB maintains for us. The second index are the indexes that we have created.

The default index for `_id` is created and painted on every collections by mongoDB automatically. This means if we are filtering for `_id` or sorting by `_id` which is then the default sort order or order by which the document are fetched, mongoDB is utilising the index for that at least.

Working with Indexes

Configuring Indexes

The `_id` index that we get out of the box for this field is actually unique by default. This is a setting MongoDB gives us to ensure that we cannot add another document with the same value in the same collection. There are use cases where we also need that behaviour for a different field and therefore we can add our own unique indexes.

For example, if we wanted email to be a unique index. We would create an index on the email field and then pass in a second argument to the `createIndex` command. The second argument allows us to configure the index – this is where we can set the `unique` option to `true`.

```
$ db.contacts.createIndex( { email: 1 }, { unique: true } )
```

If we execute this command, we may receive a duplicate key error collection if we already have have duplicate values within our collection. This will also show the document(s) where the duplicate key field exists. This is an advantage of the unique index because we would get such a warning if we try to add it or we already have it in place and tried to add a document with a value that already existed. Unique indexes can help us as developers to ensure data consistency and avoid duplicate data for fields that we need to have as unique. This index is not only useful to speed up our find queries but also to guarantee that we have unique values for that given field in that collection.

Working with Indexes

Understanding Partial Filters

Another interesting kind of configuration for a filter is setting up a so-called partial filter.

For example, if we were creating an application for calculating what someone will get once they retire – we would typically only look for a person older than 60. Having an index on the `dob.age` field might make sense. The problem of course is that we have a lot of values in our index that we never actually query for. Now the index will still be efficient but it will be unnecessarily big and an index eats up size on our disk. Additionally, the bigger the index is, the more performance certain queries will take nonetheless.

If we know certain values will not be looked at or only very rarely and we would be fine using a full collection scan, we can actually create a partial index where we only add the values we are regularly going to look at.

```
$ db.contacts.createIndex( { "dob.age": 1 }, { partialFilterExpression: { "dob.age": { $gt: 60 } } } )
```

```
$ db.contacts.createIndex( { "dob.age": 1 }, { partialFilterExpression: { gender: "male" } } )
```

We can add this option to compound indexes as well. In the `partialFilterExpression`, we define which field will narrow down the set of values we want to add (*we can use a totally different field e.g. gender*). We can use all the equality expression we have previously seen e.g. `$gt`, `$lt`, `$exist` etc.

The second expression would create a index on the age but only for elements where the underlying document is for a male while the first will only create a index on age but for elements where the underlying document is for person older than 60.

If we only created a partial index using the second example and performed the below query, we would notice that MongoDB will perform a full collection scan and will ignore the partial index. This is because MongoDB determined that yes we are looking for a field that is part of the index (age) but since we did not search for gender in our query, it considered the partial index too risky use and MongoDB as a top priority ensures that we do not lose any data. therefore, the results we receive back will also include documents that has a gender of female and not male only because it performed a full collection scan and not a partial index scan.

```
$ db.contacts.explain( ).find( { "dob.age": { $gt: 60 } } )
```

In order for MongoDB to use the partial index we must also filter by gender:

```
$ db.contacts.explain( ).find( { "dob.age": { $gt: 60 }, gender: "male" } )
```

The difference between a partial index and a compound index, for partial indexes the overall index is smaller. In the above example only the ages of males are stored and female keys are not sorted in the index and therefore the index size is smaller leading to a lower impact on our hard drive. Also our right queries are sped up because if we insert a female, that will never have to be added to the

index. This still make a lot of sense if we often filter for this type of combination i.e. for the age and then only males – a partial index can make a lot of sense if we rarely look for the other result i.e. we rarely look for women.

Whenever mongodb has the impression that our find request would yield more than what's in our index, it will not use that index but if we typically run queries where we are within our index (filtered or partial index) then mongodb will take advantage of it. We would then benefit from having a smaller index and having less impact with writes.

So again it depends on the application we're writing and whether we often just need a subset or whether we typically need to be able to query everything, in which case a partial index won't make much sense.