



{ name: mongo, type: DB }



mongoDB

MongoDB

The Complete Developer's Guide

Introduction to MongoDB

What is MongoDB?

MongoDB is a database which is created by the company who is also called MongoDB. The name stems from the word "humongous". This database is built to store a lot of data but also being able to work with the huge data efficiently. Ultimately, this is a database solution.

There are many database solutions such as MySQL, PostgreSQL, TSQL etc.

MongoDB is most importantly a database server that allows us to run different databases on it for example a Shop database. Within the database we would have different collections such as a Users collection or a Orders collection. We can have multiple databases and multiple collections per database.

Inside of the collection we have something called documents. Documents look like JavaScript JSON objects. Inside of a collection the documents are schema-less and can contain different data. This is the flexibility that MongoDB provides us with whereas SQL based database are very strict about the data stored within the database tables. Therefore, the MongoDB database can grow with the application needs. MongoDB is a No-SQL database.

Typically we will need some kind of structure in a collection because applications typically requires some type of structure to work with the data.

Diagram 1.1:



JSON (BSON) Data Format:

```
{  
  "name": "Alex",  
  "age": 29,  
  "address": {  
    "city": "Munich"  
  },  
  "hobbies": [  
    { "name": "Cooking" },  
    { "name": "Football" }  
  ]  
}
```



The above is an example of the JSON data format. A single document is surrounded by curly brackets. The data is normally structured with a Keys. Keys consist of a Name of the Key and a Key value. The Name of the Key (*which will be referred to as Key from now on*) and the Key Value must be wrapped around quotation marks (unless if the data is a type of number).

There are different types of values we can store such as: string, number, booleans and arrays.

We can also nest documents within documents. This allows us to create complex relations between

data and store them within one document, which makes working with the data and fetching data more efficient because it is contained in one document in a logical way. SQL in contrast requires more complex method of fetching data which require joins to find data in table A and data in table B to retrieve the relevant data.

Behind the scenes on the server, MongoDB converts the JSON data to a binary version of the data which can be stored and queried more efficiently. We do not need to concern ourselves with BSON as we would tend to work with JSON data.

The whole theme of MongoDB is flexibility, optimisation and usability and it is what really sets MongoDB apart from other database solutions because it is so efficient from a performance perspective as we can query data in the format we need it instead of running complex restructuring on the server.

The Key MongoDB Characteristics.

MongoDB is a no SQL solution because it is following an opposite concept/philosophy to SQL based databases. Instead of normalising the data i.e. storing data distributed across multiple tables where every table has a clear schema and then using relations, MongoDB goes for storing data together in a document. It does not force a schema hence schema-less/No-SQL.

We can have multiple documents in a single collection and they can have different structures as we have seen in Diagram 1.1. This is important, it can lead to messy data but it still our responsibility as developers to work with clean data and to implement a solution that works. On the other hand this provides us with a lot of flexibility. We could use MongoDB for applications that might still evolve, where the exact data requirements are not set yet. MongoDB allows us to started and we could always add data with more information in the same collection at a later point in time.

We also work with less relations. There are some relations, but with these embedded (nested) documents, we have less collections (tables) which we connect but instead we store data together. This is where the efficiencies is derived from, since data is stored together and when we fetch data from our application it does not require to reach out to multiple tables and merge the data because all the data is already within the single collection. This is where the speed, performance and flexibility comes from and can be seen beneficial for when building applications. This is the main reason why No-SQL solutions are so popular for read and write heavy applications.

MongoDB Ecosystem

The below Diagram 1.2 is the current snapshot of the MongoDB companies ecosystem and product offerings. The focus of this guide is on the MongoDB database used locally on our machines and on the cloud using Atlas. We will also dive in Compass and the Stitch world of MongoDB.

Diagram 1.2:



Installing MongoDB

MongoDB runs on all Operating Systems (OS) which include Windows/Mac/Linux. To install MongoDB we can visit their webpage on:

<https://www.mongodb.com/>

Under products select MongoDB server and download the MongoDB Community Server for our OS platform of choice. Install the MongoDB Server by following the installation steps.

Important Note: On Windows when installing click on the Custom Setup Type. MongoDB will be installed as a service which will be slightly different to how MongoDB runs on Mac & Linux.

On Mac and Linux we simply have a extracted folder which contains files. We would copy all the contents within this folder and paste them into any place within our OS i.e wherever we would want to install MongoDB.

We would then want to create a folder called data and a sub-folder called db anywhere within our OS, preferably in the root of the OS.

On Windows open up the command prompt or on Mac/Linux open up the terminal. This is where

we are going to spend most of our time using special commands and queries. Run the following command:

```
$ mongo
```

This should return command not found.

To fix this problem on a Mac go to the user folder and find a file called `.bash_profile` file (*if this does not exist we could simply create it*). Edit the file using a text editor. Add the following line:

```
export PATH=/Users/Username/mongobd/bin:$PATH
```

The path should be wherever we placed the MongoDB binary zip files. We need to add `:$PATH` at the end on Mac/Linux to make sure all our other commands work on our OS. Save the file and close the file.

Important Note: if you run into a problem on not being able to edit the `.bash_profile` using text editor use the following command to edit it within the terminal:

```
$ sudo nano ~/.bash_profile
```

This will allow you to edit the file within the terminal and enter the mongo bin file path. Press CTRL + o to save and CTRL + x to exit the nano edit.

To fix this on a Windows OS, we need to create an environment variable. Press the windows key and type environment which should suggest the Edit Environment Variable option. Under the user variables edit Path to add the directory path to where we installed the MongoDB files:

C:\Program Files\MongoDB\Server\4.0\bin

Restart the terminal/command prompt and now run the command:

\$ mongo

This should now return a error of connect failed on Mac/Linux.

On Windows it will connect because MongoDB is running as a service and has already started as a background service because we would have checked this during the installation. If we open the command prompt as administrator and ran the command 'net stop MongoDB' this will stop the background service running automatically and we can manually start and stop the MongoDB service running on windows. DO NOT RUN THIS COMMAND ON MAC/LINUX.

The mongo command is the client which allows us to connect to the server and then run commands on the databases. To start the service on Mac/Linux we would use the following command:

\$ mongod

When running this command to start the server it may fail if we chose a different default /data/db folder. If we used a different folder and not within the root of our OS we would need to start the mongod command instance followed by the --dbpath flag and the place where the /data/db is located if not within the root directory.

```
$ sudo mongod --dbpath "/data/db"
```

On Mac we would need to run the mongod command every time we wish to run the mongoDB service whereas on Windows this will run automatically even after restarting the system.

Now that we have the mongod server running minimise the terminal on Mac/Linux and open up a new terminal. We cannot close the mongod server terminal because it is running the service and if closed the mongoDB server everything will stop working and we cannot continue to work with the database server. Pressing the CTRL + C keys within the terminal will quit the mongod service, but we would need to re-run the mongod command again should we wish to run the server again.

We are now in the mongo shell which is the environment where we can run commands against our database server. We can create new databases, collections and documents which we will now focus on in the following sections.

Time to get Started

Now that we have the mongod server running and we can now connect to it using the mongo shell we can now enter the following basic commands in the mongo terminal:

Command	Description
<code>\$ cls</code>	Clear the terminal.
<code>\$ show dbs</code>	Display existing databases (there are three default databases: admin, config and local which store meta data).
<code>\$ use databaseName</code>	Connect/Switch to a database. If the database does not exist it will implicitly create a new database using the databaseName. It will not create the database until a collection and document is added.
<code>\$ db.collectionName.insertOne({"name of key": "key value"})</code>	Create a new collection. The db relates to the current connected database. This will implicitly create a new collection if it does not exist. We must pass at least one new data in the collection using the .insert() command passing in a JSON object. This will return the object to confirm the data was inserted into the database.

Important Note: we can omit the quotes around the name of the key within the shell but we must contain the quotes for the key value unless the key value is type of number. This is a feature within the mongo shell which work behind the scenes. MongoDB will also generate a uniqueid for new documents inserted into the collection.

Command	Description
<code>\$ db.collectionName.find()</code>	Display the document within the database collection.
<code>\$ db.collectionName.find().pretty()</code>	Display the documents within the database collection but prettify the data in a more humanly readable format.

This is a very basic introductory look at the following shell commands we can run in the mongo terminal to create a new database, switch to a database, create a collection and documents and display all the documents within a database collection either in the standard or pretty format.

Tip: to run the mongod server on a different port to the default port 27017 by run the following command. Note you would need to specify the port when running the mongo shell command as well. You would use this in case the default port is being used by something else.

```
$ sudo mongo --port 27018
```

```
$ mongo --port 27018
```

Shell vs Drivers

The shell is a great neutral ground for working with MongoDB. Drivers are packages we install for different programming languages the application might be written in. There are a whole host of drivers for the various application server languages such as PHP, node, C#, python etc. Drivers are the bridges between the programming language and the MongoDB server.

As it turns out, in these drivers, we would use the same command as we use in the shell, they are just slightly adjusted to the syntax of the language we are working with.

The drivers can be found on the MongoDB website:

<https://docs.mongodb.com/ecosystem/drivers/>

Throughout this document we will continue to use the Shell commands as it is the neutral commands. We can take the knowledge of how to insert, configure inserts, query data, filter data, sort data and many more shell commands. These commands will continue to work when we use the drivers but we would need to make reference to the driver documentation to understand how to use the shell commands but using the programming language syntax to perform the commands using the drivers. This will make us more flexible with the language we use when building applications that uses MongoDB.

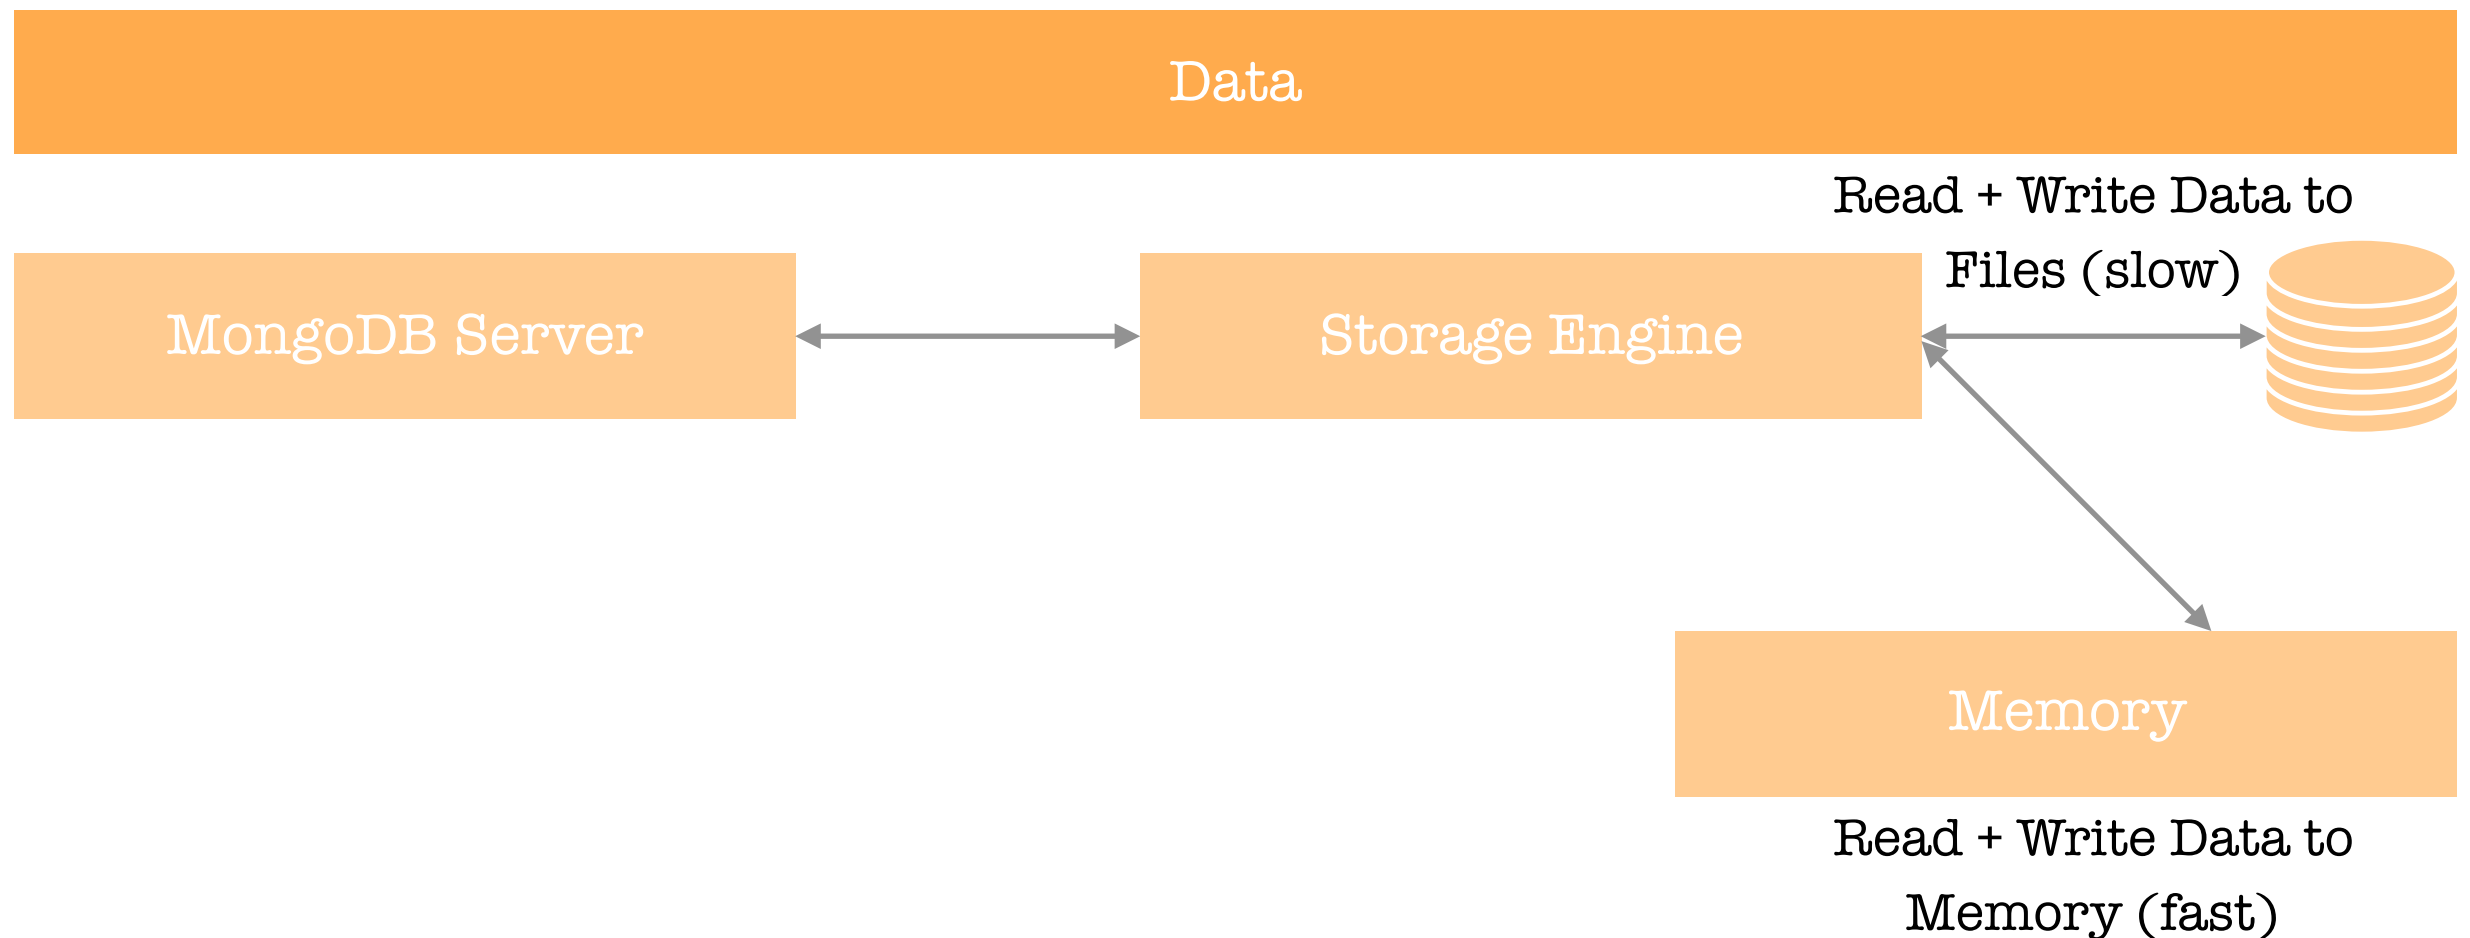
MongoDB & Clients: The Big Picture

Diagram 1.3:



MongoDB & Clients: The Big Picture

Diagram 1.4:



As we can see in Diagram 1.3 the application driver/shell communicates to the MongoDB server. The MongoDB server communicates with the storage engine. It is the Storage Engine which deals with the data passed along by the MongoDB Server, and as Diagram 1.4 depicts it will read/write to database and/or memory.

Understanding the Basics & CRUD Operations

Create, Read, Update & Delete (CRUD)

We could use MongoDB to create a variety of things such as an application, Analytics/BI Tools or data administration. In an application case, we may have an app where the user interacts with our code (the code can be written in any programming language) and the mongoDB driver will be included in the application. In the case of a Analytics/BI Tools we may use the BI Connector/Shell provided by mongoDB or another import mechanism provided by our BI tool. Finally, in the database administrator case we would interact with the mongoDB shell.

In all the above cases we would want to interact with the mongoDB server. In an application we would typically want to be able to create, read, update or delete elements e.g. a blog post app. With analytics, we would at least want to be able to read the data and as an admin we would probably want to do all the CRUD actions.

CRUD are the only actions we would want to perform with our data i.e. to create it, manage it or read it. We perform all these actions using the mongoDB server.

Diagram 1.5:

CREATE

```
insertOne(data, options)  
insertMany(data, options)
```

UPDATE

```
updateOne(filter, data, options)  
updateMany(filter, data, options)  
replaceOne(filter, data, options)
```

READ

```
findOne(filter, options)  
find(filter, options)
```

DELETE

```
deleteOne(filter, options)  
deleteMany(filter, options)
```

The above are the four CRUD operations and the commands we can run for each action. In later sections we will focus on each CRUD action individually to understand in-depth each of the actions and syntax/command we can use when performing CRUD operation with our MongoDB data collection and documents.

Understanding the Basics & CRUD Operations

Finding, Inserting, Updating & Deleting Elements

To show all the existing databases within the MongoDB server we use the command "show dbs" while we use the use followed by the database name to switch to a database. The db will then relate to the switched database.

To perform any CRUD operations, these commands must always be performed/executed on a collection where you want to create/update/delete documents. Below are example snippets of CRUD commands on a fictitious flights database (where the collection is called flightData).

```
$ db.flighData.insertOne( {distance: 1200} )
```

This will add a single document to the collection as we have seen previously.

```
$ db.flightData.deleteOne( {departureAirport: "LHR"} )
```

```
$ db.flightData.deleteMany( {departureAirport: "LHR"} )
```

The delete command takes in a filter. To add a filter we would use the curly brackets passing in

which name of key and key value of the data we wish to filter and delete. In the above example we used the `departureAirport` key and the value of `TXL`. The `deleteOne` command will find the first document in our database collection that meets the criteria and deletes it. The command will return:

```
{ "acknowledged" : true, "deleteCount" : 1 }
```

If a document was deleted in the collection this will show the number of deleted documents (the `deleteOne` command will always return 1). If no documents matched the filter and none were deleted the returned `deleteCount` value will be 0.

The `deleteMany` in contrast will delete many documents at once where the documents matches the filter criteria specified.

Note: The easiest way to delete all data in a collection is to delete the collection itself.

```
$ db.flightData.updateOne( {distance: 1200}, { $set: {marker: "delete"}} )
```

```
$ db.flightData.updateMany( {distance: 1200}, { $set: {marker: "delete"}} )
```

The `update` command takes in 3 argument/parameters. The first is the filter which is similar to the `delete` command. The second is how we would want to update/change the data. We must use the

`{ $set: { } }` keyword (anything with a \$ dollar sign in front of the keyword is a reserved word in *mongoDB*) which lets *mongoDB* know how we are describing the changes we want to make to a document. If the update key:value does not exist, this will create a new key:value property within the document else it will update the existing key:value with the new value passed in. The third parameter is options which we will analyse in great detail in the latter sections.

Important Note: when passing in a filter we can also pass in empty curly brackets `{ }` which will select all documents within the collection.

If successful with updating many this will return within the terminal an acknowledgement as seen below, where the number of matched the filter criteria and the number of data modified:

```
{ "acknowledged" : true, "matchedCount" : 2, "modified" : 2 }
```

If we were to delete all the documents within a collection and use the command to find data in that collection i.e using the `db.flightData.find().pretty()` command, the terminal will return empty/nothing as there are no existing documents to read/display.

The above demonstrates how we can find, insert, update and delete elements using the update and delete command.

Now we have seen how we can use `insertOne()` to add a single document into our collection. However, what if we want to add more than one document? We would use the `insertMany()` command instead.

```
db.flightData.insertMany( [  
  {  
    "departureAirport": "LHT",  
    "arrivalAirport": "TXL"  
  },  
  {  
    "departureAirport": "MUC",  
    "arrivalAirport": "SFO"  
  }  
])
```

We pass in an array of objects in order to add multiple documents into our database collection. The square brackets is used to declare an array. The curly brackets declare a object and we must use comma's to separate each object. If successful, this will return acknowledged of true and the `insertIds` of each object/document added into the collection.

Important Note: MongoDB by default will create a unique id for each new document which is assigned to a name of key called “_id” followed by a random generated key. When inserting a object we could assign our own unique id using the _id key followed by a unique value. If we insert a object and pass in our own _id key value and the value is not unique this will return a duplicate key error collection in the terminal. We must always use a unique id for our documents and if we do not specify a value for _id then MongoDB will generate one for us automatically.

Understanding the Basics & CRUD Operations

Diving Deeper Into Finding Data

Currently we have seen the .find() function used without passing any arguments for finding data within a collection. This will retrieve all the data within the collection. Just as we would use filter to specify a particular records or documents when deleting or updating a collection, we can also filter when finding data.

We can pass a document into the find function which will be treated as a filter as seen in the example below. This allows us to retrieve a subset of the data rather than the whole data within an application.

```
db.flightData.find( {intercontinental : true } ).pretty()
```

We can also use logical queries to retrieve more than one document within a collection that matches the criteria as demonstrated in the below example. We query using another object and then one of the special operators in MongoDB.

```
db.flightData.find( {distance: { $gt: 1000 } } ).pretty()
```

In the above we are using the \$gt: operator which is used for finding documents “greater than” the value specified. If we were to use the findOne() operator this will return the first record within the collection that matches the criteria.

Understanding the Basics & CRUD Operations

Update vs UpdateMany

Previously we have seen the updateOne() and updateMany() functions. However, we can also use another update function called update() as seen in the example below:

```
db.flightData.update( { _id: ObjectId("abc123") }, { $set: { delayed: true } } )
```

The update() function works exactly like the updateMany() function where all matching documents to the filter are updated. The difference between update() and updateMany() is that the \$set:

operator is not required for the update() function whereas this will cause an error for either the updateOne() and updateMany() functions. So we can write the above syntax like so and would not get an error:

```
db.flightData.update( { _id: ObjectId("abc123") }, { delayed: true } )
```

The second and main difference is that the update function takes the new update object and replaces the existing object (*this does not affect the unique id*) updating the document. It will only patch the update object instead of replacing the whole existing object (just like the updateOne() and updateMany() functions), if we were to use the \$set: operator, otherwise it would override the existing document.

This is something to be aware of when using the update() function. If we intend to replace the whole existing document with a new object then we can omit the \$set: operator. In general it is recommended to use updateOne() and updateMany() to avoid this issue.

If, however, we want to replace a document we should use the replaceOne() function. Again, we would place our filter and the object we want to replace with. This is a more explicit and more safer way of replacing the data in a collection.

```
db.flightData.replaceOne( { id: ObjectId("abc123") }, { departureAirport: "LHT", distance: 950 } )
```

Understanding the Basics & CRUD Operations

Understanding Find() & The Cursor Object

If we have a passengers collection which stores the name and age of passengers and we want to retrieve all the documents within the passenger collection we can use the find() function as we have seen previously.

```
db.passengers.find().pretty()
```

Useful Tip: when writing commands in the shell we can use the tab key to autocomplete for example if we wrote db.passe and tab on our keyboard, this should auto-complete db.passengers.

We will notice where a collection has many data, the find() function will return all the data but display all the data with the shell. If we scroll down to the last record we should see Type "it" for more within the shell. If we type the command it and press enter, this will display more data from the returned find() function. The find() command in general returns back what is called a Cursor Object and not all of the data.

The find() does not give an array of all the documents within a collection. This makes sense as the

collection could be really large and if the find() was to return the whole array, imagine if a collection had 2million documents – this could take a really long time but also send a lot of data over the connection.

The Cursor Object is an object that has many meta data behind it that allows us to cycle through the results, which is what the “it” command did. It used the Cursor Object to fetch the next group (cycle) of data from the collection.

We can use other methods on the find() function such as toArray() which will exhaust the cursor i.e. go through all of the cursors and fetch back all the documents within the array (i.e. not stopping after the first 20 documents – a feature within the mongoDB shell).

```
db.passengers.find().toArray()
```

There is a forEach method that can also be used on the find() function. The forEach allows us to write some code to do something on every element that is in the database. The syntax can be found within the driver documents for whichever language we are using for our application e.g. PHP or JavaScript etc. Below is a JavaScript function which the shell can also use:

```
db.passengers.find().forEach( (document) => { printjson(document) } )
```

The `forEach` function in JavaScript gets the document object passed in automatically into the arrow function and we can call this whatever we want i.e. `passengersData`, `data`, `x`, etc. In the above we called this `document`. We can then use this object and do whatever we want i.e. we used the `printjson()` command to print/output the document data as JSON. The above will also return all the documents within the collection because the `forEach` loops on every Cursor Object.

To conclude, the `find()` function does not provide us with all the documents in a collection even though it may look like it in some circumstances where there are very little data within a collection. Instead it returns a Cursor Object which we can cycle through the return more documents from the collection. It is unto us as the developer to use the cursor to either force it to get all the documents from a collection and place it in an array or better using the `forEach` or other methods to retrieve more than 20 documents (*the default number of items returned in the shell*) from the collection. Note the `forEach` is more efficient because it fetches/returns objects on demand through each iteration rather than fetching all the data in advance and loaded into memory which saves both on bandwidth and memory.

The Cursor Object is also the reason why we cannot use the `.pretty()` command on the `findOne()` function because the `findOne` returns one document and not a Cursor Object. For Insert, Update and Delete commands the Cursor Object does not exist because these methods do not fetch data, they simply manipulate the data instead.

Understanding the Basics & CRUD Operations

Understanding Projections



Imagine in our database we have the data for a person record and in within our application we do not need all the data from the document but only the name and age to display on our web application. We could fetch all the data and filter/manipulate the data within our application in any programming language. However, this approach will still have an impact on the bandwidth by fetching unnecessary data – something we want to prevent. It is better to filter the data out from the MongoDB server and this is exactly what projection allows us to do.

Below are examples of using projections to filter the necessary data to retrieve from our find query.

```
db.passengers.find( {}, {name: 1} ).pretty()
```

We need to pass in a first argument to filter the find search (note: a empty object will retrieve all documents). The second argument allows us to project. A projection is setup by passing another document but specifying which key:value pairs we want to retrieve back. The one means to include it in the data returned to us.

The above will return all the passengers document but only the name and id, omitting the age from the returned search results. The id is a special field in our data and by default it is always included. To exclude the id from the returned results, we must explicitly exclude it. To exclude something explicitly we would specify the name of key and set the value to zero as seen below:

```
db.passengers.find( {}, {name: 1, _id:0} ).pretty()
```

Note: we could do the same for age (e.g. age: 0), however, this is not required because the default is everything but the _id is not included in the projection unless explicitly specified using the one.

The data transformation/filtering is occurring on the mongoDB server before the data is shipped to us and is something that we would want because we do not want to retrieve unnecessary data which will impact on the bandwidth.

Understanding the Basics & CRUD Operations

Embedded Documents & Arrays

Embedded documents is a core feature of MongoDB. Embedded documents allows us to nest other documents within each other and having one overarching document in the collection.

There are two hard limits to nesting/embedded documents:

1. We can have up to 100 level of nesting (a hard limit) in MongoDB.
2. The overall document size has to be below 16mb

The size limit for documents may seem small but since we are only storing text and not files (we would use file storage for files), 16mb is more than enough.

Along with embedded documents, another documents we can store are arrays and this is not strictly linked to embedded documents, we can have arrays of embedded documents, but arrays can hold any data. This means we have list of data in a document.

Below are examples of embedded documents and arrays.

```
db.flightData.updateMany( {}, {$set: {status: {description: "on-time", lastUpdated: "1 hour ago"}}} )
```

In the above example we have added a new document property called status which has a embedded/nested document of description and lastUpdated. If we output the document using .find() function, the below document would now look something like the below:

```
{
  "_id": ...,
  "departure": "LHT",
  "arrivalAirport": "TXL",
  "status": {
    "description": "on-time",
    "lastUpdated": "1 hour ago",
  }
}
```

Note: we could add more nested child documents i.e. description could have a child nested document called details and that child could have further nested child documents and so on.


```
db.passengers.updateOne( {name: "Albert Twostone"}, {$set: {hobbies: ["Cooking", "Reading"]} } )
```

Arrays are marked with square brackets. Inside the array we can have any data, this could be multiple documents (i.e. using the curly brackets {}), numbers, strings, booleans etc.

If we were to output the document using the .find() function, the document would look like something below:

```
{
  "_id": ...,
  "name": "Albert Twostone",
  "age": 63,
  "hobbies": [
    "Cooking",
    "Reading"
  ]
}
```

Albert Twostone will be the only person with hobbies and this will be a list of data. It is important to note that hobbies is not a nested/embedded document but simply a list of data.

Understanding the Basics & CRUD Operations

Accessing Structured Data

To access structured data within a document we could use the following syntax:

```
db.passengers.findOne( {name: "Albert Twostone"} ).hobbies
```

We can specify the name of a structured data within a document by using the find query and then using the name of key we wish to access from the document, in the above we wanted to access the hobbies data which will return the hobbies array as the output:

```
["Cooking", "Reading"]
```

We can also search for all documents that have hobbies of Cooking using the syntax below as we have seen previously. This will return the whole document entry where someone has Cooking as a hobby. MongoDB is clever enough to look in arrays to find documents that match the criteria.

```
db.passengers( {hobbies: "Cooking"} ).pretty()
```

Below is an example of searching for objects (this includes searching within nested documents):

```
db.flightData.find( {"status.description": "on-time"} ).pretty()
```

We use the dot notation to drill into our embedded documents to query our data. It is important that we wrap the dot notation in quotations (e.g. "status.description") otherwise the find() function would fail.

This would return all documents (the whole document) where the drilled criteria matches. This allows us to query by nested/embedded documents. We can drill as far as we need to using the dot notation as seen in the example below:

```
db.flightData.find( {"status.details.responsible": "John Doe"} ).pretty()
```

This dot notation is a very important syntax to understand as we would use this a lot to query our data within our mongoDB database.

Understanding the Basics & CRUD Operations

Conclusion

We have now covered all the basic and core features of mongoDB to understand how mongoDB works and how we can work with it i.e. store, update, delete and read data within the database as well as how we can structure our data.

Understanding the Basics & CRUD Operations

Resetting The Database

To purge all the data within our MongoDB database server we would use the following command:

```
use databaseName
```

```
db.dropDatabase()
```

We must first switch to the database using the use command followed by the database name. Once we have switched to the desired database we can reference the current database using db and then call on the dropDatabase() command which will purge the specified database and its data.

Similarly, we could get rid of a single collection in a database using the following command:

```
db.myCollection.drop()
```

The myCollection should relate to the collection name.

These commands will allow us to clean our database server by removing the database/collections that we do not want to keep on our MongoDB server.

Schemas & Relations: How to Structure Documents

Why Do We Use Schemas?

There is one important question to ask – wasn't MongoDB all about having **no** data Schemas i.e. Schema-less. To answer this question, MongoDB enforces no Schemas. Documents do not have to use the same schema inside of one collection. Our documents can look like whatever we want it to look like and we can have totally different documents in one and the same collection i.e. we can mix different schemas.

Schemas are the structure of one document i.e. how does it look like, which fields does it have and what types of value do these fields have. MongoDB does not enforce schemas; however, that does not mean that we cannot use some kind of schema and in reality we would indeed have some form of schema for our documents. It is in our interest if we were to build a backend database that we have some form of structure to the types of documents we are storing. This would make it easier for us to query our database and get the relevant data and then cycle through this data using a programming language to display the relevant data within our application.

We are most likely to have some form of schemas because we as developers would want it and our applications will need it. Whilst we are not forced to have a schema we would probably end up with some kind of schema structure and this is important to understand.

Schemas & Relations: How to Structure Documents

Structuring Documents



We can use any of the structured approach in the diagram above depending on how we require it in our applications. In reality we would tend to use the approach in the middle or on the right.

The middle approach used the best of both worlds where there are some structure to the data, however, it also has the flexibility advantage that MongoDB provides us so that we can store extra information.

Note: we can assign the null value to properties in order to have a structured approach although the data may not have any actual values associated with the property. A null value is considered a valid value and therefore we can use a SQL (structured) type approach with all our documents.

There is no single best practice with how to set the structure of our data within our documents and it is up to us as developers to use the best structure that works best for our applications or whichever is to our personal preference.

Schemas & Relations: How to Structure Documents

Data Types

Now that we understand that we are free to define our own schemas/structure for our documents, we are now going to analyse the different data types we can use in MongoDB.

Data Types are the types of data we can save in the fields within our documents. The below table break the different data types for us:

Type	Example Value
String	"John Doe"
Boolean	TRUE
NumberInt (int32)	55, 100, 145
NumberLong (int64)	100000000000
NumberDecimal	12.99
ObjectId	ObjectId("123abc")
ISODate	ISODate("2019-02-09")
Timestamp	Timestamp(11421532)
Embedded Documents	{"a": {...}}
Arrays	{"b": [...]}

Notice how the text type requires quotation marks (*single or double*) around the value. There are no limitation in the size of the text. The only limitation is the 16mb for the whole document. The larger the text the larger the data it takes.

Notice how numbers and booleans do not require a quotation marks around the value.

There are different types of numbers in mongoDB. Integer (int32) are 32bit long numbers and if we try to store a number longer than this they would overflow that range and we will end up with a

different number. For longer integer numbers we would use NumberLong (int64). The integer solution we decide to choose will dictate how much space will be allocated and eaten up by the data. Finally, we can also store NumberDecimal i.e. numbers with decimal values (a.k.a float in other programming languages).

The default within the shell is to store a int64 floating point value but we also have a special type of NumberDecimal provided by MongoDB to store high precision floating point values. Normal floating point values (a.k.a doubles) are rounded and are not precise after their decimal place. However, for many use cases the floating point (double) is enough precision required e.g. shop store. If we are performing scientific calculations or something that requires a high precision calculation, we are able to use the special type that offers this very high decimal place precision (i.e. 34 digits after the decimal place).

The ObjectId is a special value that is automatically generated by MongoDB to provide a unique id but it also provides some temporal component that allows for sorting built into the ObjectId, respecting a timestamp.

The above table provides all the data types within MongoDB that we can use to store data within our database server.

Schemas & Relations: How to Structure Documents

Data Types & Limits

MongoDB has a couple of hard limits. The most important limitation: a single document in a collection (including all embedded documents it might have) must be less than or equal to 16mb. Additionally we may only have 100 levels of embedded documents.

We can read more on all the limitation (in great detail) on the below link:

<https://docs.mongodb.com/manual/reference/limits/>

For all the data types that mongoDB supports, we can find a detailed overview on the following link:

<https://docs.mongodb.com/manual/reference/bson-types/>

Important data type limits are:

- Normal Integers (int32) can hold a maximum value of +-2,147,483,674
- Long Integers (int64) can hold a maximum value of +-9,223,372,036,854,775,807
- Text can be as long as we want – the limit is the 16mb restriction for the overall document.

It's also important to understand the difference between int32 (NumberInt), int64 (NumberLong) and a normal number as you can enter it in the shell. The same goes for a normal double and NumberDecimal.

NumberInt creates a int32 value => NumberInt(55)

NumberLong creates a int64 value => NumberLong(7489729384792)

If we just use a number within the shell for example `insertOne({a: 1})`, this will get added as a normal double into the database. The reason for this is because the shell is based on JavaScript which only knows float/double values and does not differ between integers and floats.

NumberDecimal creates a high precision double value => NumberDecimal("12.99")

This can be helpful for cases where we need (many) exact decimal places for calculations.

When working with MongoDB drivers for our application's programming language (e.g. PHP, .NET, Node.js, Python, etc.), we can use the driver to create these specific numbers. We should always browse the API documents for the driver we are using within our applications to identify the methods for building int32, int64 etc.

Finally we can use the `db.stats()` command in the MongoDB shell to see stats of our database.

Schemas & Relations: How to Structure Documents

How to Derive Our Data Structure Requirements

Below are some guidelines to keep to mind when we think about how to structure our data:

- What data does our App need to generate? What is the business model?

User Information, Products Information, Orders etc. This will help define the fields we would need (and how they relate).

- Where do I need my data?

For example, if building a website do we need the data on the welcome page, products list page, orders page etc. This help define our required collections and field groupings.

- Which kind of data or information do we want to display?

For example the welcome page displays product names. This will help define which queries we need i.e. do we need a list of products or a single product.

These queries we plan to have also have an impact on our collections and document structure.

MongoDB embraces the idea of planning our data structure based on the way we retrieve the data so that we do not have to perform complex joins but we retrieve the data in the format or almost in

the format we need it in our application.

- How often do we fetch the data?

Do we fetch data on every page reload, every second or not that often? This will help define whether we should optimise for easy fetching of data.

- How often do we write or change the data?

Do we change or write data often or rarely will help define whether we should optimise for easy writing of data.

The above are things to keep in mind or to think about when structuring our data structures and schemas.

Schemas & Relations: How to Structure Documents

Understanding Relations

Typically we would have multiple collections for example a users collection, a product collection and a orders collections. If we have multiple collections that are relatable or where the documents in these relations are related, we obviously have to think about how do we store related data.

Do we use embedded documents because this is one way of reflecting a relation or alternatively, do we use references within our documents?

Nested/ Embedded Documents

Customers Collection

```
{
  "userName": "John",
  "age": 28,
  "address": {
    "street": "First Street",
    "City": "Chicago"
  }
}
```

References

Customers Collection:

```
{
  "userName": "Alan"
  "favBooks": ["id1", "id2"]
}
```

Books Collection:

```
{
  "_id": "id1",
  "name": "Lord of the Rings"
}
```

In the reference example above, we would have to run two queries to join the data from the different collections. However, if a book was to change, we would only update it in the books collection as the id would remain the same whereas in a embedded document relation we would have to update multiple customer records affected with the new change.

Schemas & Relations: How to Structure Documents

One to One Embedded Relation Example

Example:

One patient has one disease summary, a disease summary belongs to one patient.



Code snippet:

```
$ use hospital  
$ db.patients.insertOne( { name: "John Doe", age: "25", diseaseSummary: { diseases: ["cold",  
"sickness"] } } )
```

Where there is a strong one to one relation between two data, it is ideal to use a one to one embedded approach as demonstrated in the above example.

The advantage of the embedded nested approach is that within our application we only require a single find query to fetch the necessary data for the patient and disease data from our database collection.

Schemas & Relations: How to Structure Documents

One to One Reference Relation Example

Example:

One person has one car, a car belongs to one person.



Code snippet:

```
$ use.carData
```

```
$ db.persons.insertOne( { name: "John", age: 30, salary: 30000 } )
```

```
$ db.cars.insertOne( { model: "BMW", price: 25000, owner: ObjectId("5b98d4654d01c") } )
```

In most one to one relationships we would generally use the embedded document relations. However, we can opt to use a reference relation approach as we are not forced to use one approach.

For example, we have a more analytics use case rather than a web application and we have a use case where we are interested in analysing the person data and or analysing our car data but not so much in a relation. In this example we have a application driven reason for splitting the data.

Schemas & Relations: How to Structure Documents

One to Many Embedded Relation Example



Example:

One question thread has many answers, one answer belongs to one question thread.

Code snippet:


```
$ use support
```

```
$ db.questionThreads.insertOne( { creator: "John", question: "How old are you?", answers: [ { text: "I am 30." }, { text: "Same here." } ] } )
```

A scenario where we may use a embedded one to many relation would be post and comments. This is because you would often need to fetch the question along with the answers in an application perspective. Also usually there are not too many answers to worry about the 16mb document limit.

Schemas & Relations: How to Structure Documents

One to Many Reference Relation Example



Example:

One city has many citizens, one citizen belongs to one city.

Code snippet:

```
$ use cityData
```

```
$ db.cities.insertOne( { name: "New York City", coordinates: { lat: 2121, lng: 5233 } } )
```

```
$ db.citizens.insertMany( [ { name: "John Doe", cityId: ObjectId("5b98d6b44d") }, { name: "Bella Lorenz", cityId: ObjectId("5b98d6b44d") } ] )
```

In the above scenario we may have a database containing a collection of all major cities in the world and a list of every single person living within that city. It would seem to make sense to have a one to many embedded relationship, however, from an application prospective we may wish to only retrieve the city data only. Furthermore, a city like New York may have over 1 million people data and this would make fetching the data slow due to the volume of data passing through the wire. Furthermore, we may end up running into the document size limit of 16mb. In this type of scenario, it would make sense to split the data up and using the reference relation to link the data.

In the above we would only store the city metadata and will not store any citizen reference as this will also end up being a huge list of citizens unique id. Instead, we would create a citizens collection and within the citizens data we would make reference to the city reference. The reference can be anything but must be unique ie. we could use the ObjectId() or the city name etc.

This will ensure that we do not exceed the limitation of the 16mb per document as well as not retrieving unnecessary data if we are only interested in returning just the cities metadata from a collection.

Schemas & Relations: How to Structure Documents

Many to Many Embedded Relation Example



Example:

One customer has many products (via orders), a product belongs to many customers.

Code snippet:

```
$ use shop
```

```
$ db.products.insertOne( { title: "A Book", price: 12.99 } )
```

```
$ db.customers.insertOne( { name: "Cathy", age: 18, orders: [ { title: "A Book", price: 12.99, quantity: 2 } ] } )
```

We would normally model many to many relationships using references. However, it is possible to

use the embedded approach as seen above. We could store a collection for the products as meta data for an application to retrieve the data in order to help populate the embedded document of the customer collection using a programming language.

A disadvantage to the embedded approach is data duplication because we have the title and price of the product within the orders array as the customer can order the product multiple times as well as other customers which will cause a lot of duplication.

If we decide to change the data for the product, not only do we need to change it within the product collection but we also have to change it on all the orders affected by this change (or do we actually need to change old orders?). If we do not care about the product title changing and the price changing i.e. we have an application that takes a snapshot of the data, we may not worry too much about duplicating that data because we might not need to change it in all the places where we have the duplicated the data if the original data changes – this highly depends on the application we build. Therefore a embedded approach may work.

In other case scenarios where we absolutely need the latest data everywhere, a reference approach may be most appropriate in a many to many relationship. It is important to think about how we would fetch our data and how often do we want to change it and if we need to change it everywhere or are duplicate data fine before deciding which approach to adopt for many to many.

Schemas & Relations: How to Structure Documents

Many to Many Reference Relation Example



Example:

One book has many authors, an author belongs to many books.

Code snippet:

```
$ use bookRegistry  
  
$ db.books.insert( { name: "favourite book", authors: [ objectId("5b98d9e4"), objectId("5b98d9a7") ]  
  })  
  
$ db.authors.insertMany( [ { name: "Martin", age: 42 }, { name: "Robert", age: 56 } ] )
```

The above is an example of a many to many relation where a reference approach may be suitable for a scenario where the data that changes needs to be reflected everywhere else.

Schemas & Relations: How to Structure Documents

Summarising Relations

We have now explored the different relation options that are available to use. This should provide us enough knowledge to think about relations and when to use the most appropriate approach depending on:

- the application needs
- how often data changes
- if a snapshot data suffice
- how large is the data (how much data do we have).

Nested/Embedded Documents – group data together logically. This makes it easier when fetching the data. This approach is great for data that belong together and is not overlapping with other data. We should always avoid super-deep nesting (100+ levels) or extremely long arrays (16mb size limit per document).

References – split data across collections. This approach is great for related data but also shared data as well as for data which is used in relations and standalone. This allows us to overcome nesting and size limits (by creating new documents).

Schemas & Relations: How to Structure Documents

Using \$lookup for Merging Reference Relations

MongoDB has a useful operation called \$lookup that allows us to merge related documents that are split up using the reference approach.

The image on the right provides a scenario of a reference approach where the customer and books have been split into two



collections. The lookup operator is used as seen below. This uses the aggregate method which we have not currently learned.

```
$ customer.aggregate( [
  { $lookup: { from: "books", localField: "favBooks", foreignField: "_id", as: "favBookData" } }
])
```

The \$lookup operator allows us to fetch two related documents merged together in one document within one step (rather than having to perform two steps). This mitigates some of the disadvantages of splitting our documents across multiple collections because we can merge them in one go.

This uses the aggregate method framework (which we will dive into in later chapters) and within the aggregate we pass in an array because we can define multiple steps on how to aggregate the data. For now we are only interested in one step (a step is a document we pass into an array) where we pass the \$lookup step. The lookup passes in a document as a value, where we define 4 attributes:

- **from** – which other collection do we want to relate documents i.e. we would pass in the name of the collection where the other document lives that we wish to merge.
- **localField** – in the collection we are running the aggregate function on, where can the reference to the other (from) collection be found in i.e. the key that stores the reference.
- **foreignField** – which field are we relating to in our target collection (i.e. the from collection)
- **as** – provide an alias for the merged data. This will become the new key which the merged data will sit.

This is not an excuse to always using a reference relation approach because this costs more performance than having an embedded document.

If we have a references or want to use a references, we have the lookup step in the aggregate method that we can use to help get the data we need. This is a first look at aggregate and we will explore what else the aggregate can do for us in later chapters.

Schemas & Relations: How to Structure Documents

Understanding Schema Validation

MongoDB is very flexible i.e. we can have totally different schemas and documents in one and the same collection and that flexibility is a huge benefit. However, there are times where we would want to lock down this flexibility and require a strict schema.

Schema validation allows mongoDB to validate the incoming data based on the schema that we have defined and will either accept the incoming data for the write or update to the database or it will reject the incoming data and the database is not changed by the new data and the user gets an error.

validationLevel

Which document get validated?

Strict

All inserts & updates

Moderate

All inserts & updates
to correct documents

validationAction

What happens if validation fails?

Error

Throw error and deny
insert/update

Warn

Log warning but
proceed

Schemas & Relations: How to Structure Documents

Adding Collection Document Validation

To add schema validation in MongoDB and the easiest method is to add validation when we create a new collection for the very first time explicitly (not implicitly when we add a new data). We can use the `createCollection` to create and configure a new collection:

```
$ db.createCollection("posts", { validator: { $jsonSchema: { bsonType: "object", required: ["title",  
"text", "creator", "comments"], properties: { title: { bsonType: "string", description: "must be a string  
and is required." }, text: { bsonType: "string", description: "must be a string and is required" },  
creator: { bsonType: "objectId", description: must be an objectId and is required }, comments:  
{ bsonType: "array", description: "must be an array and is required", items: { bsonType: "object",  
required: ["text,"], properties: { text: { bsonType: "string", description: "must be a string and is  
required" }, author: { bsonType: "objectId", description: "must be an objectId and is required" }}}}} }  
}} )
```

The first argument to the `createCollection` method is the name of the collection i.e. we are defining the name of the collection. The second argument is a document where we would configure the new collection. The validator is an important piece of the configuration.

The validator key takes in another sub document where we can now define a schema against incoming data where inserts and updates has to validated. We do this by inserting a \$jsonSchema key with another nested sub document which will hold the schema.

We can add a son type with the value of object, so that everything that gets added to the collection should be a valid document or object. We can set a required key which has an array value. In this array we can define names of fields in the document which will be part of the collection that are absolutely required and if we try to add data that does not have these fields, we will get an error or warning depending on our settings.

We can add a properties key which is another nested document where we can define how for every property of every document that gets added to the collection will look like. In the example above we defined the title property, which is a required property, in more detail. We can set the bsonType which is the data type i.e. string, number, boolean, object, array etc. We can also set a description for the data property.

Because an array and has multiple items, we can add an items key and describe how the items should look like. We can nest this and this can have another nested required and properties keys for the items objects that exists within the array.

So the Keys to remember are:

The bsonType key is the data type.

The required key is an array of required properties that must be within an insert/update document.

The properties key defines the properties. This has sub key:value of of bsonType and description.

The Item key defines the array items. This can have sub key:value of all the above.

Important Note: it may be difficult to read in the terminal and may be easier to write in a text editor first and then paste into the terminal to execute the command. We can call the file validation.js to save the collection validation configuration. Visual Studio/Atom/Sublime or any other text editor/IDE will help with auto-formatting. Visual Studio has a option under code > Preference > Keyboard Shortcuts and then you can search for a shortcut command such as format document (shortcut is Shift + Option + F on a Mac).

We can now validate the incoming data when we explicitly create the new collection. We can copy the command from the text editor and paste it back into the shell and run the command to create the new collection with all our validation setup. This will return `{ "OK" : 1 }` in the shell if the new collection is successfully created.

If a new insert/update document fails the validation rules, the new document will not be added to the collection.

Schemas & Relations: How to Structure Documents

Changing the Validation Action

As a database administrator we can run the following command:

```
$ db.runCommand( { collMod: "post", validator: {...}, validationAction: "warn" } )
```

This allows us to run administrative commands in the shell. We pass a document with information about the command we wish to run. For example, in the above we run a command called collMod which stands for collection modifier, whereby we pass in the collection name and then we can pass in the validator along with the whole schema.

We can amend the validator as we like i.e. add or remove validations. In the above we added another administrative command after the validator document as a sibling called validationAction. The validationLevel controls whether all inserts and updates are checked or only updates to elements which were valid before. The validationAction on the other hand will either throw an "error" and stop the insert/update action or "warn" of the error but allow the insert/update to occur. The warn would have written a warning into our log file and the log file is stored on our system. We can update the validation action later using the runCommand() method as seen above.

Schemas & Relations: How to Structure Documents

Conclusion

Things to consider when modelling and structuring our Data.

- **In which format will we fetch your data?**

How does the application or data scientists need the data? We want to store the data in a way that it is easy to fetch especially in a use case where we would fetch a lot.

- **How often will we fetch and change the data?**

Do we need to optimise for writes or reads? It is often for reads but it may be different depending on the scenario. If we write a lot then we want to avoid duplicates. If we read a lot then maybe some duplicates are OK, provided these duplicates do not change often.

- **How much data will we save (and how big is it)?**

If the data is huge, maybe embedding is not the best choice.

- **How is the data related (one to one, one to many, many to many)?**

- **Will duplicate data hurt us (=> many Updates)?**

Do we update our data a lot in which we have to update a lot of duplicates. Do we have snapshot data where we do not care about updates to the most recent data.

- **Will we hit the MongoDB data/storage limit (embed 100 level deep and 16mb per document)?**

Modelling Schemas

- Schemas should be modelled based on application needs.
- Important factors are: read and write frequencies, relations, amount (and size) of data.

Schema Validation

- We can define rules to validate inserts and updates before writing to the database.
- Choose the validation level and action based on the application requirements.

Modelling Relations

- Two options: embedded documents or references.
- Use embedded documents if we have one-to-one or one-to-many relationships and there are no app or data size reasons to split the data.
- Use reference if data amount/size or app needs require it or for many-to-many relations.
- Exceptions are always possible — keep the app requirements in mind!

Useful Articles & Documents:

<https://docs.mongodb.com/manual/reference/limits/>

<https://docs.mongodb.com/manual/reference/bson-types/>

<https://docs.mongodb.com/manual/core/schema-validation/>

Exploring The Shell & The Server

Setting dbpath & logpath

In the terminal we can run the following command to see all the available options for our MongoDB server:

```
$ mongo --help
```

This command will provide a list of all the available options we can use to setup/configure our MongoDB server. For example the `--quiet` option allows us to change the way things get logged or output by the server.

Note: use the official document on the MongoDB website for more detailed explanation of all the available options.

The `--dbpath` arg and `--logpath` arg allows us to configure where the data and log files gets stored to because MongoDB writes our data to real files on our system. The logs allows us to see for example warnings of json schema validation as we seen in the last section.

We can create folders such as `db` and `logs` (these can be named as anything we want) and have

these folders located anywhere we want for example we could create it within the MongoDB directory which contains the bin folder and other related files.

If we start using mongod instance without any additional settings, it will use the root folder that has a data/db folder to store all our database records as a default setting. However, we can use the settings above to tell mongod to use another folder directory to store our data, the same is true for our logs.

When we start the instance of our MongoDB server, we can run the following command and passing in the options to declare the path of the dbpath and logpath as seen below:

Mac/Linux:

```
$ sudo mongod --dbpath /Users/userName/mongoDB/db
```

Windows command:

```
$ mongod --dbpath \Users\userName\mongoDB\db
```

Enter our password and this should bring up our MongoDB server as we have seen previously. We should now see in the db folder, MongoDB has created a bunch of files as it is now saving the data

in the specified folder that we passed into our command. This is now using a totally different database storage for writing all our data which is detached from the previous database storage of the default database path. Running the following command will also work for our logs:

Mac/Linux:

```
$ sudo mongod --dbpath /Users/userName/mongoDB/db --logpath /Users/userName/mongoDB/logs/logs.log
```

Windows command:

```
$ mongod --dbpath /Users/userName/mongoDB/db --logpath  
\Users\userName\mongoDB\logs\logs.log
```

The logs folder path requires a log file which we would define with a .log extension. This will automatically create and add a logs.log file within the directory path if the file does not exist when we run the command. All the output in the terminal will now be logged in the logs.log file compared to previously where it was logged in the terminal shell. This file can be reviewed for persistent and auditing of our server and viewing any warnings/errors.

This is how we set custom paths for our database and log files.

Exploring The Shell & The Server

Exploring the MongoDB Options

If we explore the different options in MongoDB using the `mongod --help` command in the terminal, there are many setup options available to us.

The WiredTiger options is related to our storage engine and we could either use the default settings or change some configurations if we know what we are doing.

We have useful commands such as `--repair` which we could run if we have any issues connecting or any warnings or issues related to our database files being corrupted. We could use the command `--directoryperdb` which will store each database in its own separate directory folder.

We could change the storage engine using the `--storageEngine` arg command, which by default is set to WiredTiger. Theoretically, MongoDB supports a variety of storage engines but WiredTiger is the default high performance storage engine. Unless, we know what we are doing and have a strong reason to change the engine, we should stick to the default.

There are other settings in regards to security which we will touch in the latter chapters.

Exploring The Shell & The Server

MongoDB as a Background Service

In the mongoDB options, there is an option called `--fork` which can only run on Mac and Linux.

```
$ mongod --fork --logpath /Users/userName/mongoDB/logs/logs.log/
```

The above fork command will error if we do not pass in a logpath to the log file. This command will start the mongoDB server as a child process. This does not block the terminal and we can continue to type in other commands in the same terminal with the server running. The server is now running as a background process instead of a foreground process which usually blocks the terminal window. In other words the mongoDB server is now running as a service (a service in the background). Therefore, in the same terminal we could run the mongo command to connect to the background mongoDB server service. This is also the reason why we require to pass in a logpath because the service is running in the background and it cannot log error/warning messages in the terminal, instead it will use/write the warning and errors in the log file.

On Windows, the fork command is unavailable. However, on Windows we can still startup mongoDB server as a service if we checked this option at the installation process. If we right click on command prompt and run as administrator, we can run the following command:

```
$ net start MongoDB
```

This will start up the MongoDB server as a background service. The question then becomes, how do we stop such a service?

On Mac we can stop the service by connecting to the server with the mongo shell and then switching to the admin database and running the shutdown server command to shut down the server we are connected to. Example commands below:

```
$ use admin
```

```
$ db.shutdownServer()
```

The exact same approach as the above will work on Windows. On Windows we also have an alternative method by opening the command prompt as administrator and running the following command:

```
$ net stop MongoDB
```

This is how we can use MongoDB server as a background service (instead of a foreground service) on either Mac, Linux or Windows.

Exploring The Shell & The Server

Using a Config File

Now that we have seen the various options we can set and use to run our MongoDB server, it is also worth noting that we can save our settings in a configuration file.

<https://docs.mongodb.com/manual/reference/configuration-options/>

This file could be automatically created for us when we run our MongoDB server, else we could create the config file ourselves and save this anywhere we want. We could create the config file within the Users/username/MongoDB/bin folder using a text editor such as VS Code to add the configuration code:

storage:

dbPath: "/Users/username/mongoDB/db/"

systemLog:

destination: file

path: "/Users/username/mongoDB/logs/logs.log/"

We can look at the documents or google search for more comprehensive config file setup.

Once we have the config file setup, how do we use the config file when we run an instance of the MongoDB server? MongoDB does not automatically pickup this file when we start to run the MongoDB server, instead when starting MongoDB we can use the following command to specify the config file the server should use:

```
$ sudo mongod --config /Users/userName/mongoDB/bin/mongod.cfg
```

```
$ sudo mongod -f /Users/userName/mongoDB/bin/mongod.cfg
```

Either above command will prompt MongoDB to use the config file from the path specified. This will start the MongoDB server with the settings setup in the configuration file. This is a useful feature because it allows us to save a snapshot of our settings (reusable blueprint) in a separate file which we can always use when starting up our MongoDB server. This also saves us time on writing a very long command prompt with all our settings when starting up our MongoDB server each time.

Important Note: we could use either .cfg or .conf as the file extension name when creating the MongoDB configuration file.

Exploring The Shell & The Server

Shell Options & Help

In this section we will go over the various shell options available to for us to use. Similar to the MongoDB server, there is a help option for the MongoDB shell:

```
$ mongo --help
```

This will provide all the command options for the shell. This has less options compared to the server because the shell is just a connecting client at the end of the day and not a server. We can use the shell without connecting to a database (if we just want to run javascript code) using the `--nodb` command, or we could use the `--quiet` command to have less output information in the terminal, we can define the port and host for the server using the `--port` arg and `--host` arg commands (by default it uses `localhost:27017`) and many more other options.

We can also add Authentication Options informations which we will learn more in later chapters.

In the shell we also have another command we can run:

```
$ help
```

This command will output a shortlist of some important help information/commands we can execute in the shell. We can also dive deeper into the help by running the help command followed by the command we want further help on, for example:

```
$ help admin
```

This will show further useful commands that we can execute when using the admin command e.g. admin.hostname() or admin.pwd() etc.

We can also have help displayed for a given database or collection in a database. For example:

```
$ use test
```

```
$ db.help()
```

We would now see all the commands that we did not see before that we can use on the new “test” database. We can also get help on the collection level which will provide a list of all the commands we can execute at the collection level.

```
$ db.testCollection.help()
```

Useful Links:

<https://docs.mongodb.com/manual/reference/configuration-options/>

<https://docs.mongodb.com/manual/reference/program/mongo/>

<https://docs.mongodb.com/manual/reference/program/mongod/>

Using the MongoDB Compass to Explore Data Visually

Exploring MongoDB Compass

We can download MongoDB Compass from the below link:

<https://www.mongodb.com/products/compass>

This is a GUI tool to interact with our MongoDB database. Once downloaded and installed on our machines we are ready to use the GUI tool. It is important to have the mongod server running in the background when we open the MongoDB Compass to connect to the database. We would connect to a Host and this by default will have localhost and port 27017. We can click connect and this will connect the GUI tool to the mongod server. We should be able to see the 3 default databases of admin, config and local.

We can now use the GUI tool to create a new database and collection name. Once a database and collection has been created we can then insert documents to the collection. We can also query our database documents.

We can now start using a GUI tool to interact with our database, collections and documents.

Note: it is best practice to learn how to use the shell first before using GUI tools.

Diving Into Create Operation

Understanding insert() Methods

We already understand that there are two methods for inserting documents into MongoDB which are `insertOne()` and `insertMany()` as an alternative. The most important thing to note is that `insertOne()` takes in a single document and we can but do not need to specify an id because we will get one automatically. The `insertMany()` does the same but with an array (list) of documents.

There is also a third alternative method for inserting documents called `insert()` – below is an example:

```
$ db.collectionName.insert()
```

This command is more flexible because it takes both a single document or an array of documents. `Insert` was used in the past but `insertOne` and `insertMany` was introduced on purpose so that we are more clear about what we will be inserting. Previously, in application code it was difficult to tell with the `insert` command whether the application was inserting a single or multiple documents and therefore may have been error prone.

There is also an importing data command as seen below:

```
$ mongoimport -d cars -c carsList --drop --jsonArray
```

The insert method can still be used in mongoDB but it is not recommended. The insert() method works with both a single document and multiple documents as seen in the examples below:

```
$ db.persons.insert( { name: "Annie", age: 20 } )
```

```
$ db.persons.insert( [ { name: "Barbara", age: 45 }, { name: "Carl", age: 65 } ] )
```

The output message in the terminal is also slightly different i.e. we would receive a text of:

```
$ WriteResult( { "nInserted" : 1 } )
```

```
$ BulkWriteResult( { "writeErrors": [], "writeConcernErrors": [], "nInserted": 2, "nUpserted": 0, "nMatched": 0, "nModified": 0, "nRemoved": 0, "upserted": [] } )
```

The above does not mean that the inserted document did not get an autogenerated id. The insert method will automatically create an ObjectId but will not display the ObjectId unlike the insertOne and insertMany commands output messages which does display the ObjectId. We can see the advantages of insertOne and InsertMany as the output message is a little more meaningful/helpful as we can immediately work with the document using the ObjectId provided (i.e. we do not need to query the database to get the new document id).

Diving Into Create Operation

Working With Ordered Inserts

When inserting documents we can define or specify some additional information. Lets look at an example of a hobbies collection where we keep track of all the hobbies people could possibly have when we insert many hobbies. Each hobby is a document with the name of the hobby:

```
$ db.hobbies.insertMany( [ { _id: "sports", "name": "Sports" }, { _id: "cooking", "name": "Cooking" },  
{ _id: "cars", "name": "Cars" } ] )
```

The id's for these hobbies can be auto-generated. However, there may be times when we want to use our own id because the data may have been fetched from some other database where we already have an existing id associated or maybe we need a shorter id. We can use `_id` and assign a value for the id. In the above the hobby name could act as a good id because each hobby will be unique. We must use `_id` and not just `id` if we want to set our own id for our documents.

Furthermore, the id must be unique else this would not work. We will no longer see an `ObjectId()` for these documents as we have used the `_id` as the unique identifier for the documents inserted.

If we try to insert a document with the same id we would receive an error message in the terminal referencing the index number (mongoDB uses zero indexing) of the document that failed the insert operation along with a description of duplicate key error.

```
$ db.hobbies.insertMany( [ { _id: "yoga", "name": "Yoga" }, { _id: "cooking", "name": "Cooking" }, { _id: "hiking", "name": "Hiking" } ] )
```

The above would fail due to the duplicate key error of cooking which was inserted previously in the above command. However, we would notice on the first item in the insertMany array i.e. Yogo will be inserted into the hobbies collectio, but the cooking and hiking documents will not be inserted into the collection due to the error. This is the default behaviour of mongoDB and this is called an ordered insert.

An ordered insert simply means that every element we insert is processed standalone, but if one fails, it cancels the entire insert operation but does not rollback the elements it has already inserted. This is important to note because it cancels the operation and does not continue to the next document (element i.e. hoking) which we would have known that it would have succeeded insert.

Often we would want this default behaviour, however, sometimes we do not. In these cases, we could override the behaviour. We would pass in a second argument, separated by a comma, to the insertMany command which is a document. This is a document that configures the insertMany operation.

```
$ db.hobbies.insertMany( [ { _id: "yoga", "name": "Yoga" }, { _id: "cooking", "name": "Cooking" }, { _id: "hiking", "name": "Hiking" } ], { ordered: false } )
```

The ordered option allows us to specify whether MongoDB should perform an ordered insert which is the default (*we could set this ordered option to true which is redundant because this is the default option*) or we could set this option to false which will make the insert operation not an ordered insert i.e. an unordered insert.

If we hit enter, we would still get a list of all the error, however, it will continue to the next document to perform the insert operation and this would insert the document that does not have any issues of duplicate keys i.e. hiking will now be inserted into the hobbies collection (yoga and cooking would fail due to the duplicate key issue).

By setting the ordered to false, we have changed the default behaviour and it is up to us to decide what we require or want in our application. It is important to note that this will not rollback the entire insert operation if something failed. This is something we will cover in the Transactions chapter. We can control whether the operation continues with the other documents and tries to insert everything that is perfectly fine.

We may use an unordered insert where we do not have much control with what is inserted into the database but we do not care about any document that fail because they already exist in the database. We could add everything that is not in the database.

Diving Into Create Operation

Understanding the writeConcern

There is a second option we can specify on `insertOne` and `insertMany` which is the `writeConcern` option. We have a client (either the shell or the application using a MongoDB server) and we have our MongoDB server. If we wanted to insert one document in our MongoDB server, on the MongoDB server we have a so called storage engine which is responsible for really writing our data onto the disk and also for managing it in memory. So our write might first end up in memory and there it manages the data which it needs to access with high frequency because memory is faster than working with the disk. The write is also scheduled to then end up on the disk, so it will eventually store data on the disk. This is true for all write operations i.e. `insertMany` and `update`.

We can configure a so-called `writeConcern` on all the write operations with an additional argument, the `writeConcern` which is another document where we can set settings.



The `w:` (default) option tells the MongoDB server of how many instances we want the write to be acknowledged. The `j:` option stands for journal which is an additional file which the storage engine manages, which is like a To-Do file. The journal can be kept to then for example perform save operations that the storage engine needs to do but have not been completed yet.

The storage engine is aware of the write and that it needs to store the data on disk just by having the write being acknowledged and being in memory. The idea behind a journal file is to make the storage engine aware of this and if the server should go down for some reason or anything else should happen, the journal file is there. If we restart the server or if the server recovers, the server can look to this file and see what it needs to do. This is a nice backup because the memory might have been wiped by then. The journal acts as a backup to-do list for the storage engine.

Writing into the database files is more performance heavy whereas a journal is like a single line which describes the write operations. Writing into the database is of course a more complex task because we need to find the correct position to insert the data and if we have indexes we also need to update these as well and therefore takes longer to perform. Writing in a to-do type list is much quicker.

We can set the `j: true` as an option which will now report a success for a write operation when it has been acknowledged and has been saved to the journal. This will provide a greater security.

There is a third option to the writeConcern which is the wtimeout: option. This simply sets the timeframe that we give out server to report a success for the write before we cancel it. For example, if we have some issues with the server connection or anything of that nature, we may simply timeout.

If we set the timeout value to a very low number, we may get more fails even though there is no actual problem, just some small latency.

```
{ w:1, j: undefined }
```

```
{ w:1, j: true }
```

```
{ w:1, timeout: 200, j: true }
```

This is the writeConcern option we can add to our write operations and how we can control this using the above document settings. Enabling the journal would mean that our writes will take longer because we do not only tell the server about the write operation but we also need to wait for the server to store the write operation in the journal, however, we get higher security that the write also succeeded. These option will again depend on our application needs.

Diving Into Create Operation

The writeConcern in Practice

Below is an example of using the writeConcern:

```
$ db.persons.insertOne( { name: "Alan", age: 44 }, { writeConcern: { w: 1, j: true } } )
```

The w:1 (default) simply means to make sure the server acknowledged the write operation. Note we could set this value to 0 which will return {"acknowledged": false} in the terminal when we insert the document. This option sends the request and immediately returns without waiting for a response of the request from the server. The storage engine had no chance of storing it in memory and to generate an objectId, hence why we receive {"acknowledged": false} in the terminal. This makes the write super fast because we do not have to wait for any response but we do not know where the write succeeded or not.

The journal by default is set to undefined or false. We can set this option to j: true. The output in the terminal does not change. The write will be slightly slower (*note if playing around locally we would not notice any change in speed*) because the engine would add the write to the journal and we would have to wait for that to finish before the operation is completed. This will provide a higher security by ensuring the write appears in the to-do list of the storage engine which will eventually

lead to the write operation occurring on the database.

Finally, the `wtimeout:` option is used to set a timeout operation. This allows us to set a time frame for the write operation so in the case where within a certain period within the year we would have shaky connections, we would rather have the write operation fail and we recognise it in our client application (we would have access to the error) and therefore try again at a later time without having to wait unnecessarily for the write operation where we have shaky connections.

Diving Into Create Operation

What is Atomicity?

Atomicity is a very important concept to any write operation. Most of the time the write operation e.g. `InsertOne()` would succeed, however, it can fail (there can be an error). These are errors that occur whilst the document is being inserted/written to memory (i.e. whilst being handled by the storage engine). For example, we were writing a document for a person including name, age and an array of hobbies, the name and age were written but then the server had issues and was not able to write the hobbies to memory. MongoDB protects us against this as it guarantees us an atomic transaction. This means the transaction either succeeds as a whole or it fails as a whole. If it fails during the write, everything is rolled back for the document we inserted.

This is important as it is on a per document level. The document means the top level document, so it includes all embedded documents and all arrays.

If we insertMany() where there are multiple documents being inserted into the database and the server fails during a write, we do not get atomicity because it only works at a document level. If we have multiple documents in one operation like the insertMany() operation, then only each document on its own is guaranteed to either fail or succeed but not on insertMany. Therefore, if we have issues during the insertMany operation, only the documents that failed are not inserted and then the exact behaviour will depend on whether we used ordered or unordered inserts but the document already inserted will not be rolled back.

We are able to control this on a bulk insert or bulk update level using a concept called transactions which we will look at in a later section as it requires some additional knowledge about MongoDB and how the service works.



MongoDB CRUD Operations are Atomic on the document level (including Embedded documents)

Diving Into Create Operation

Importing Data

To import data into our database, we must first exit the shell by pressing the control + c keys on our keyboard.

In the normal terminal, we need to navigate to a folder that contains the JSON file that we would want to import (JSON files can be imported) using the cd command. We can use the ls command to view the list of items within the directory we are currently in.

Once navigated to the folder containing the import file, we can run the following command:

```
$ mongoimport tv-shows.json -d moviesData -c movies --jsonArray --drop
```

The mongoimport command should be globally available since we added the path to our mongo binary to our path variables on our operating systems. If we did not do this, we need to navigate into the folder where our MongoDB binaries are in order to execute the mongoimport command above.

The first argument we pass is the name of the file we want to import (if we are not in the path of the located file we would have to specify the full folder path along with the file name). We then specify the database we want to import the data into using the -d flag. We can also specify the collection by using the -c flag.

If the JSON file holds array of documents we must also specify the --jsonArray flag to make the mongoimport command aware of this fact about the import data.

The last argument option we can add to the import command is the --drop which will tell the mongoimport that should this collection should already exist, it will be dropped and then re-added, otherwise it will append the data to the existing collection.

Important Note: the mongod server should be running in the background when we use the import command. When we press enter to execute the command, this will return in the terminal the connected to: localhost, dropping: moviesData.movies and imported: # documents in the terminal as a response to inform us which MongoDB server it is connected to, whether a collection was dropped/deleted from the database collection and the total number of data imported into the database collection.

Diving Into Read Operation

Methods, Filters & Operators

In the shell, we access the database with the db command (this will differ slightly in a mongoDB drive). We would get access to a database and then to a collection in the database. Now we can execute a method like find, insert, update or delete on the collection. We would pass some data into the method as parameters/arguments for the method. These are usually a key:value pair where one is the field and the other is the value for that field name (documents are all about field and values or key and values).



The argument in the above example happens to also be a filter because the find method accepts a filter. It can use a filter to narrow down the set of documents it returns to us. In the above we have a equality or single value filter where the data is exactly the criteria i.e. equality.

We can also use more complex filters as seen in the below example. We have a document which has a field and its value is another document which has an operator as a field followed by a value.



We can recognise operators by the dollar sign \$ at the beginning of the operator. These are all reserved fields which are understood by mongodb. The operator in the above example is called a range filter because it does not just filter for equality, instead this will look for all documents that have an age that is greater than (\$gt) the value i.e. 30.

this is how the Read operator works and we will look at various different operators and the different ways of using them and the different ways of filtering the data that is returned to us. This is the structure we should familiarise ourselves with for all of our Read operations.

Diving Into Read Operation

Operators and Overview

There are different operators that we can differentiate into two groups:

- Query Selectors
- Projection Operators

Query selectors such as `$gt` allows us to narrow down the set of documents we retrieve while projection operators allows us to transform/change the data we get back to some extent. Both the Query and Projection operators are Read related operators.

Aggregation allows us to read from a database but also perform more complex transforms. This concept allows us to setup pipeline of stages to funnel our data through and we have a few operators that allow us to shape the data we get back to the form we need in our application.

- Pipeline Stages
- Pipeline Operators

The Update has operators for the fields and arrays. Inserts have no operators and deletes uses the same operators as the Read operators.

How do operators impact our data?

Type	Purpose	Change Data?	Example
Query Operator	Locate Data	No	\$eq \$gt
Project Operator	Modify data Presentation	No	\$
Update Operator	Modify & Add additional data	Yes	\$inc

Diving Into Read Operation

Query Selectors & Projection Operators

There are a couple of categories for Query Selectors:

Comparison, Logical, Element, Evaluation, Array, Comments & Geospatial.

Projections Operators we have:

\$, \$elemMatch, \$meta & \$slice

Diving Into Read Operation

Understanding findOne() and find()

The findOne() method finds exactly one document. We are able to pass in a filter into the method to define which one document to return back. This will find the first matching document.

```
$ db.movies.findOne( )
```

```
$ db.movies.findOne( { } )
```

Both of the above syntax will return the first document within the database collection. Note, this does not return a cursor as it only returns one document.

The alternative to findOne() is the find() method. The find() method will return back a cursor. This method theoretically returns one document, but since it provides us a cursor, it does not give us all the document but the first 20 documents within the shell.

```
$ db.movies.find( )
```

To narrow the find search we would need to provide a filter. To provide a filter we would pass in a document as the first argument (this is true for both find and findOne methods). The difference would be that findOne will return the first document that meets the criteria while the find method

will return all documents that meet the criteria.

```
$ db.movies.findOne( { name: "The Last Ship" } )
```

```
$ db.movies.findOne( { runtime: 60 } )
```

```
$ db.movies.find( { name: "The Last Ship" } )
```

```
$ db.movies.find( { runtime: 60 } )
```

To filter the data, we would specify the name of the field/key followed by the value we are expecting to filter the field by. In the above example we are filtering the name of the movie to be "The Last Ship". By default MongoDB will try to find the filter by equality.

This is the difference between find and findOne and how we would pass in a filter to narrow down the return read results. It is important to note that there are way more operators and ways to filtering our queries to narrow down our Read results when using either of the find commands.

Diving Into Read Operation

Working with Comparison Operators

In the official documentation we can view all the various operations available to us:

<https://docs.mongodb.com/manual/reference/operator/query/>

We will explore some of the comparative operators in the below examples:

```
$ db.movies.find( { runtime: 60 } ).pretty( )
```

```
$ db.movies.find( { runtime: { $eq: 60 } } ).pretty( )
```

The \$eq operator is the exact same as the default equality query which will find the document that matches equally to the query value which in the above case is runtime = 60.

```
$ db.movies.find( { runtime: { $ne: 60 } } ).pretty( )
```

This will return all documents that have a runtime not equal to 60.

```
$ db.movies.find( { runtime: { $gt: 40 } } ).pretty( )
```

```
$ db.movies.find( { runtime: { $gte: 40 } } ).pretty( )
```

The \$gt return all documents that have a runtime greater than to 40 while the \$gte returns greater than or equal to 40.

```
$ db.movies.find( { runtime: { $lt: 40 } } ).pretty( )
```

```
$ db.movies.find( { runtime: { $lte: 40 } } ).pretty( )
```

The \$lt return all documents that have a runtime less than to 40 while the \$lte returns less than or equal to 40.

Diving Into Read Operation

Querying Embedded Fields & Arrays

We are not limited to querying top level fields and are also able to query embedded fields and arrays. To query embedded fields and arrays is quite simple as demonstrated below:

```
$ db.movies.find( { "rating.average": { $gt: 7 } } ).pretty( )
```

We specify the path to the field that we are interested in querying the data. We must put the path within quotations marks because we use the dot (which will invalidate the syntax) to detail each embedded field within the path that leads to the field we are interested in. The above example is a single level embedded document, if we wrote e.g. rating.total.average, this is a 2 level embedded document. We can make the path as deep as we need it to be and we are not limited to one level.

We can also query arrays as seen below:

```
$ db.movies.find( { genres: "Drama" } ).pretty( )
```

```
$ db.movies.find( { genres: ["Drama"] } ).pretty( )
```

The casing is important in the query. This will return all genres that has Drama included in it. Equality in an array does not mean that Drama is the only item within the array; it means that Drama exists within the array. If we wanted an exactly only Drama within the array we would use square brackets. We can also use dot to go down embed level paths that has an array.
